

# **Introduction**

## **COL331/COL633**

**Abhilash Jindal**

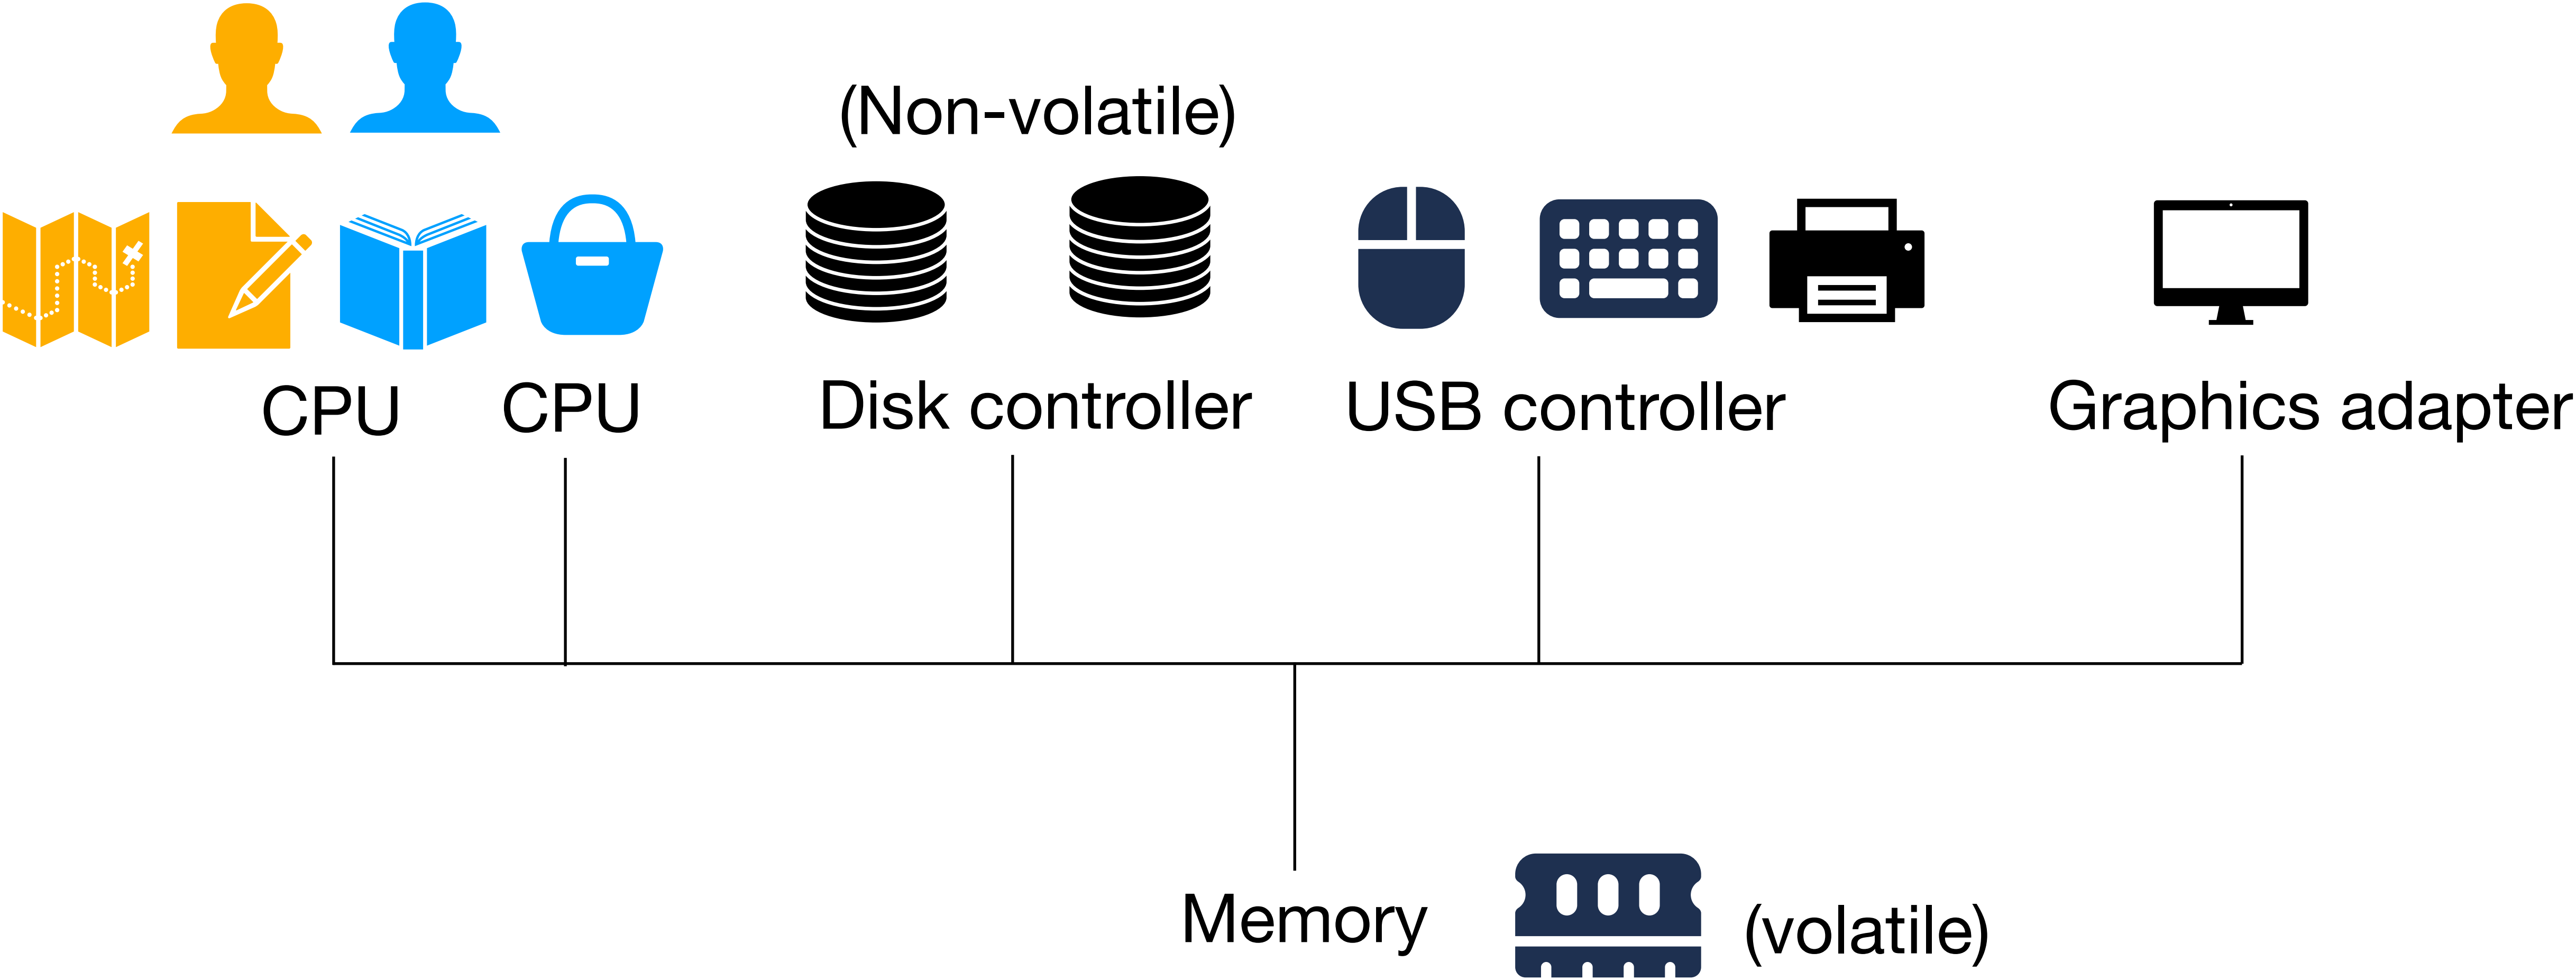
**Reference. OSTEP book: Chapter 2**

# Administrivia

- <http://abhilash-jindal.com/teaching/2023-2-col-331/>
- Grading criteria, TAs, late policy, audit criteria, quizzes, labs, project, piazza link
- Piazza access code: col331col633

**Why does OS matter to a computer?**


# Computer organization



# Purpose of an OS

- **Resource management**
- Provide higher-level services
- Protection and isolation

# Purpose of OS: Resource management



- Example: `cpu.c`
- Give the illusion of more CPUs than there are
- Multiplex the hardware

# Calculator analogy: Computing long sum



20
10
30
<b>50</b>
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

# Sharing the calculator



20
10
30
50
30
10
20
10



10
70
20
40
20
10
50
10

- Steps to share the calculator:
  - $20 + 10 = 30 + 30 = 60$
  - Write 60 in notebook, remember that we were done till 30, give calculator
  - $10 + 70 = 80$
  - Write 80 in notebook, remember that we were done till 70, give the calculator back



# Remember whatever is on the screen and give calculator?



20
10
30
50
30
10
20
10



10
70
20
40
20
10
50
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

Can I give calculator here?

“Save 1 and remember pointer to be at 10”

No! Sum would be wrong!

“+ xx = (move pointer)” has to be atomic



# Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer”
- OS switches CPU between processes in the same manner as our calculator example
- What should happen when multiple processes want to run simultaneously?
  - Fairness: One banker got more calculator time than others
    - Often need to break away from fairness. Game should get more CPU time than Dropbox to provide good user experience
  - Starvation freedom: When there are multiple bankers, one banker never got the calculator

# Purpose of OS: Resource manager

*Different approaches to memory management:*

- Segment different memory portions to different processes
- Multiplex memory pages across processes




Operating system

Computer hardware





Or



# Memory management


- Segmentation is cheap to implement
- But not flexible. What if a process needs more memory than what OS gave?
- Paging is complicated to implement
- Highly flexible



# Purpose of an OS


- Resource management
- **Provide higher-level services**
- Protection and isolation

# Purpose of OS: Provide higher-level services



Operating system

Computer hardware




*Example: io.c*

Disk interface: List of blocks



File system OS interface: Folders and files

```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```




# Why file system?

## Why not just multiplex disk blocks like memory?



Operating system


Computer hardware



- Disk blocks live after programs exits, computer restarts
- Different programs read / write same file
  - vim writes io.c
  - gcc reads io.c, write io
  - We finally run io

1	2	3	4	5	6	7	8	9	10	11	12	..	..	..
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

# Higher level services provide portability



- Abstract away hardware details
- Programs need not be rewritten when moving from hard-disk drive to solid state drive

```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```

The diagram shows a sequence of disk sectors represented as a table. The sectors are numbered 1 through 12, followed by three sectors containing double dots (..) to indicate continuation. Two arrows originate from the `write` function call in the code above, pointing to sectors 3 and 6, illustrating how data is written to specific physical locations on the disk.

1	2	3	4	5	6	7	8	9	10	11	12	..	..	..
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----




# Purpose of an OS

- Resource management
- Provide higher-level services
- **Protection and isolation**

# Purpose of OS: Protection

- S\_IRUSR | S\_IWUSR: File can only be rw by user
- Disallow inappropriate accesses. Example: users reading each other's files




```
int fd = open("/home/abhilash/exam.txt", S_IRUSR | S_IWUSR);  
int rc = write(fd, "Q1: \n", 4);  
close(fd);
```

1	2	3	4	5	6	7	8	9	10	11	12	..	..	..
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

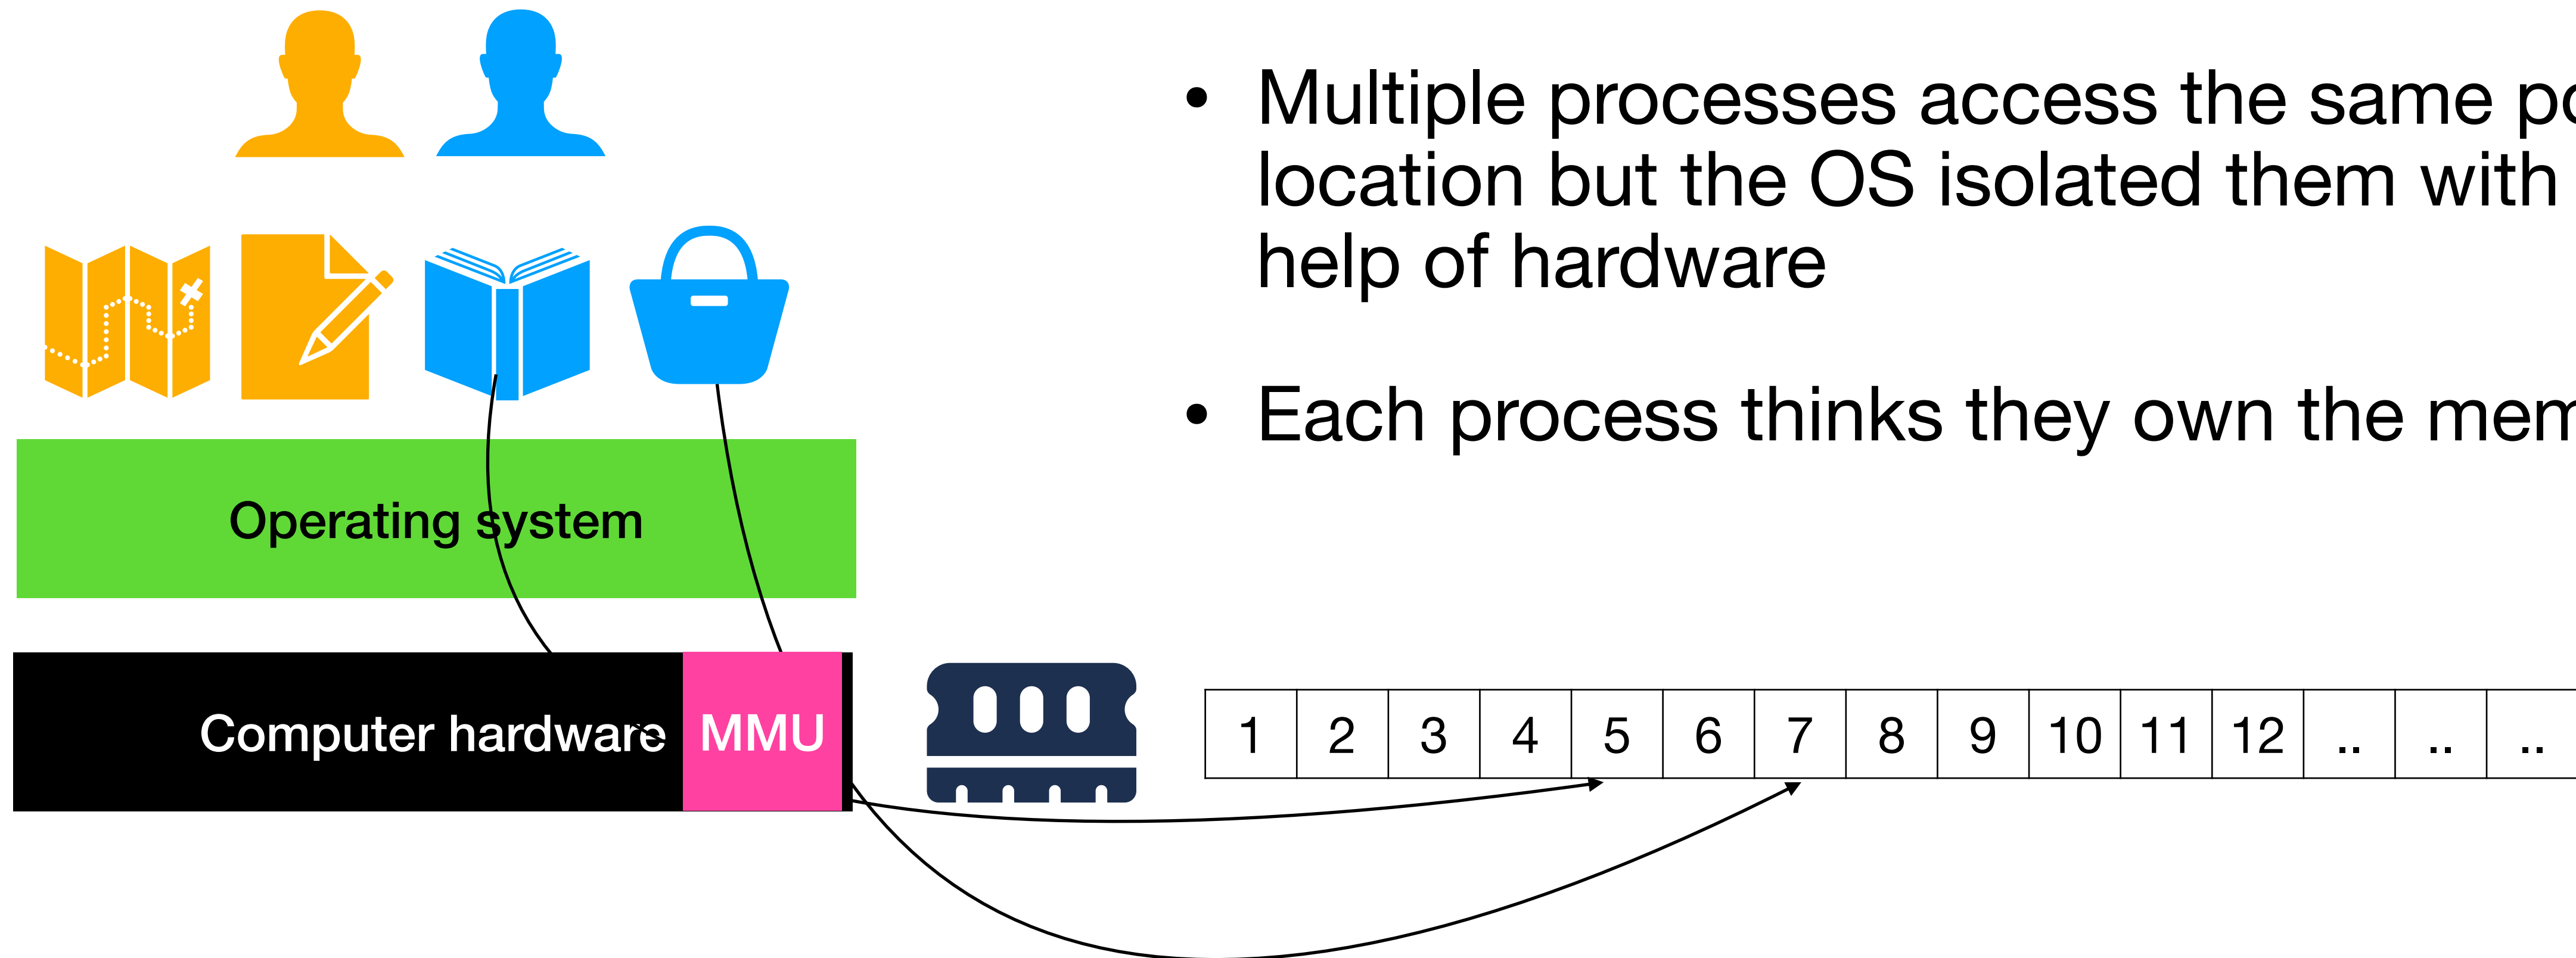
# Purpose of OS: Protection

- Disallow inappropriate accesses.  
Example: processes reading each other's memory



# Purpose of OS: Isolation

- Example: `mem.c`
- Multiple processes access the same pointer location but the OS isolated them with the help of hardware
- Each process thinks they own the memory



# Course structure: OS in action

- We will build an OS (xv6) from scratch
  - Booting: Bootloader, ELF format
  - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
  - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
  - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call
  - Concurrency: data races, different types of locks
  - Memory virtualization: memory hierarchy, address translation mechanism, demand paging, thrashing, fork system call
  - Shell: Pipes, IO redirection
  - Parallelism: Enable more CPUs, revisit locks

# Data races due to concurrency

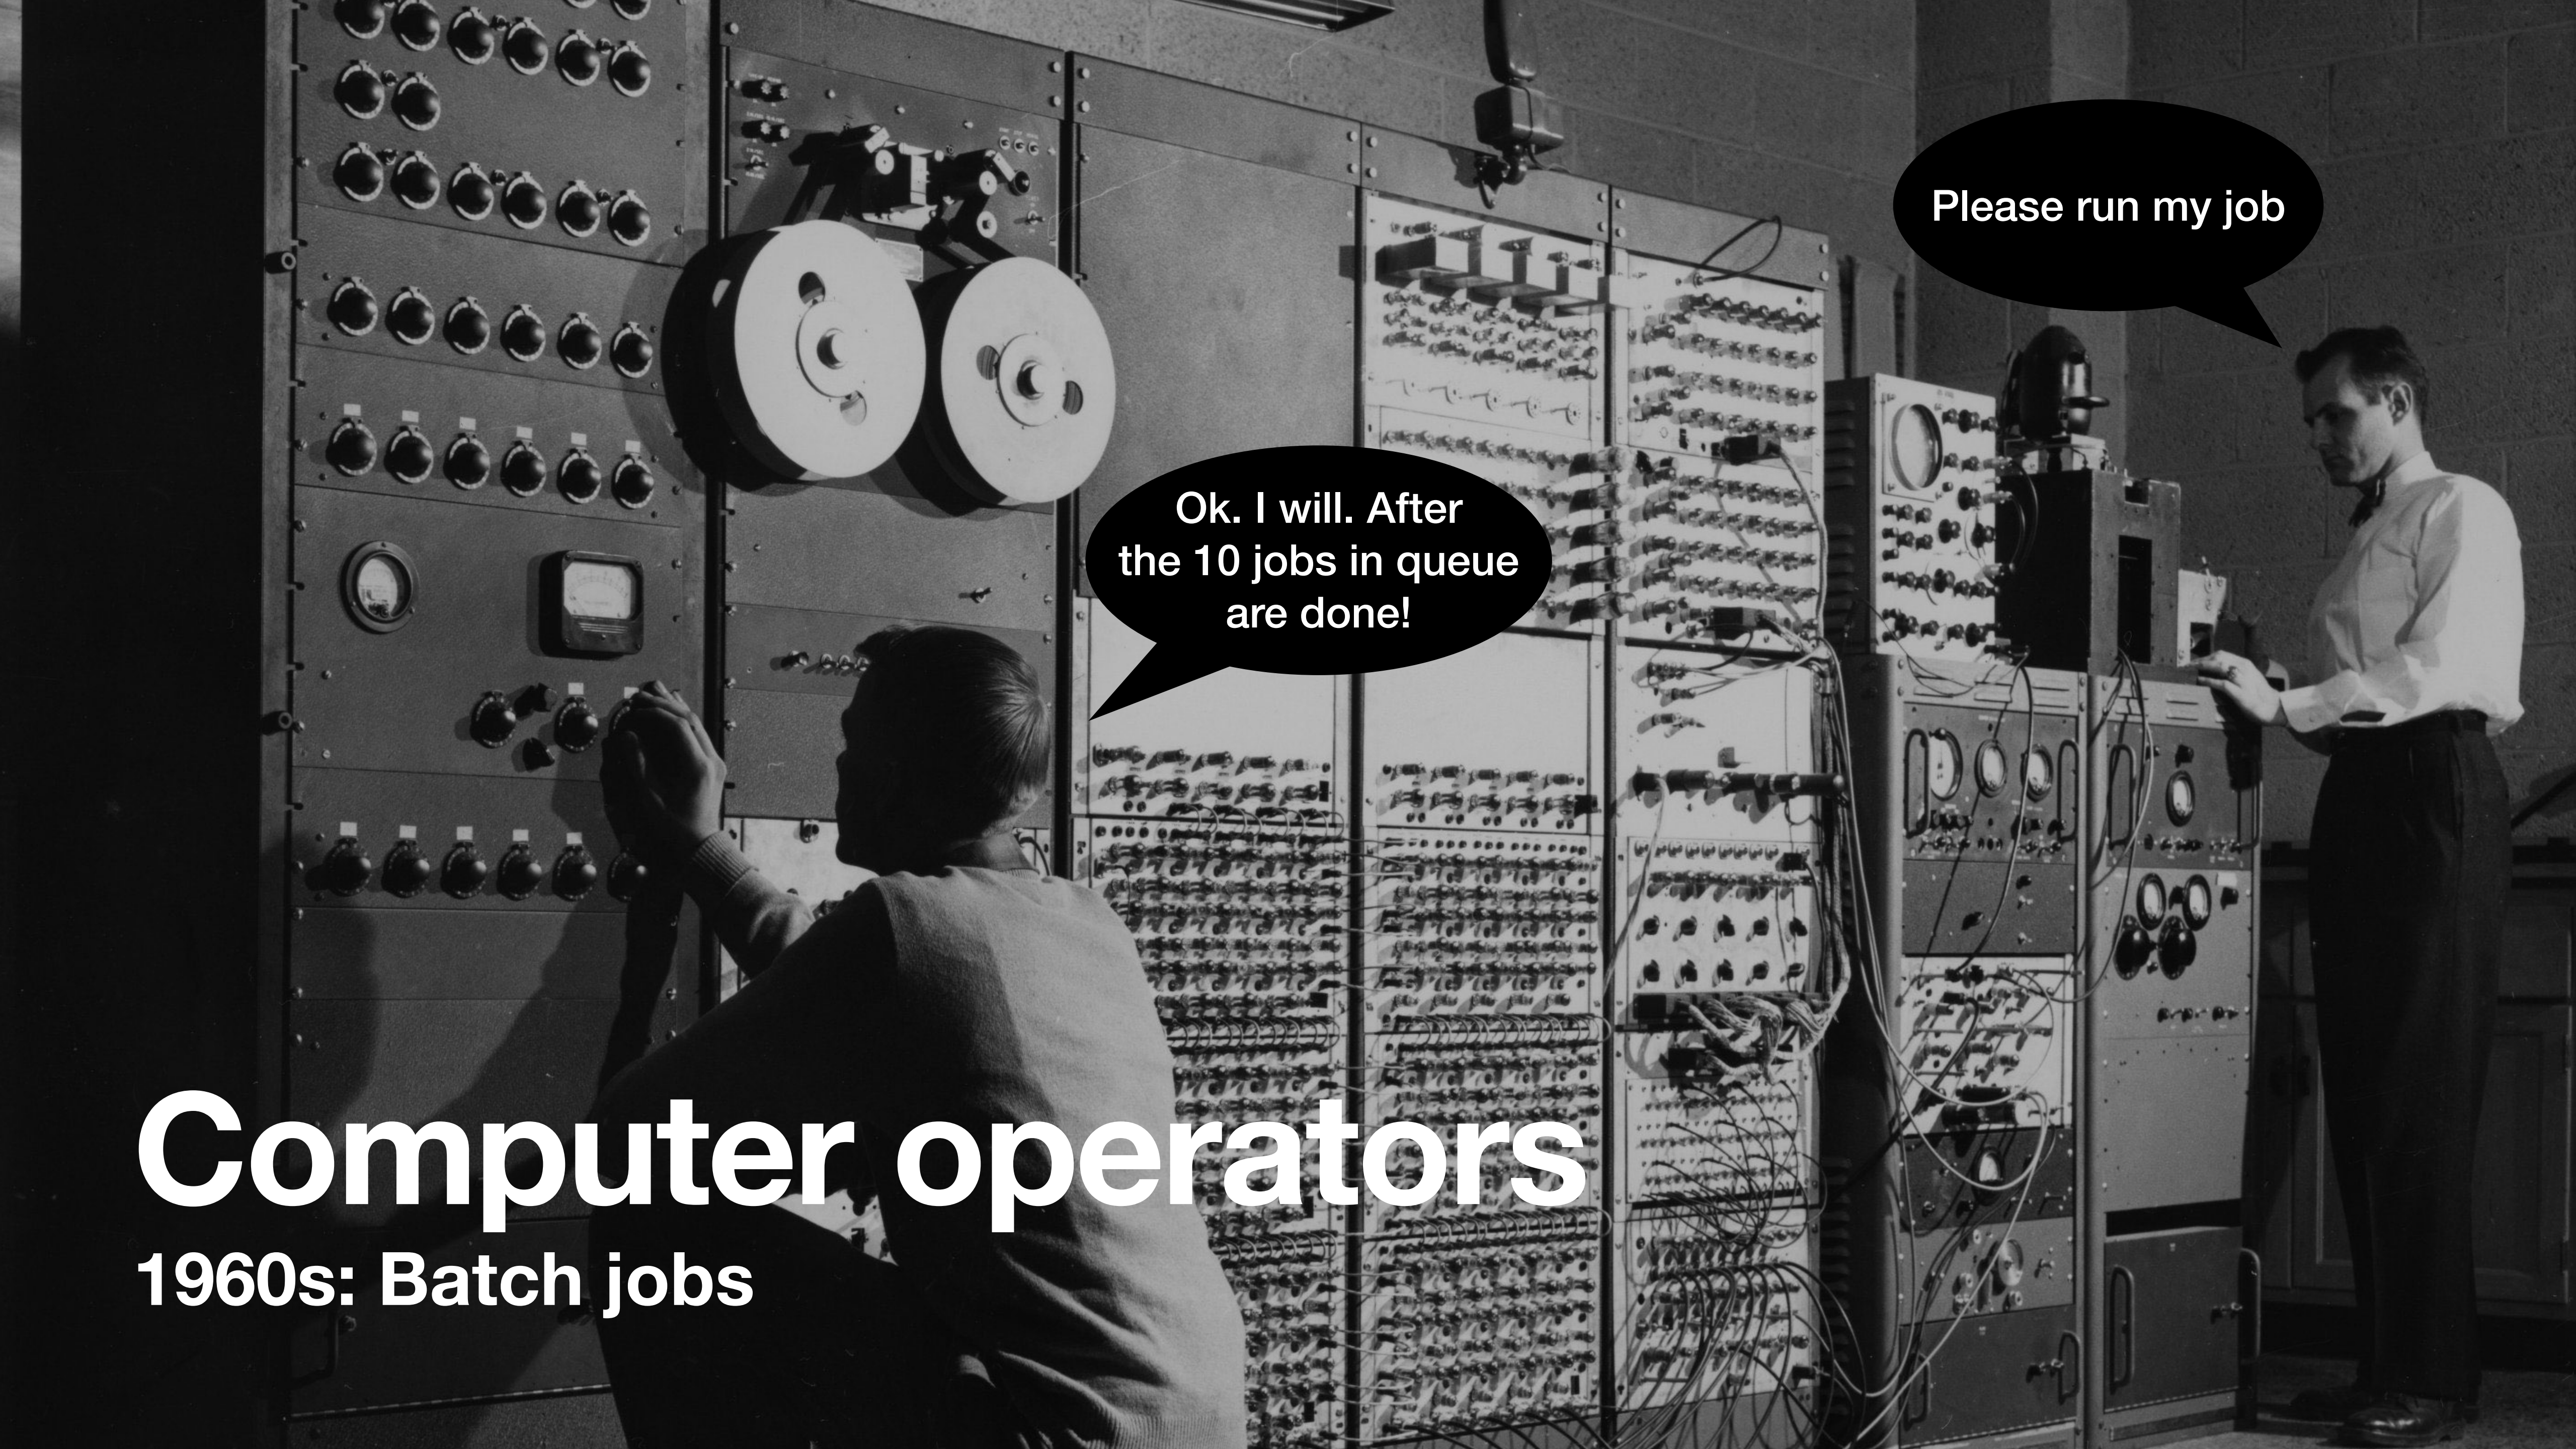
`./threads 100000`

Thread 1	Thread 2
Read counter = 0	
Write counter = 1	
	Read counter = 1
	Writer counter = 2
Read counter = 2	
	Read counter = 2
Writer counter = 3	
	Writer counter = 3

**Why should I learn OS in 2024?  
Isn't it a solved problem?**

**We have indeed made good progress ..**





Please run my job

Ok. I will. After  
the 10 jobs in queue  
are done!




# Computer operators

1960s: Batch jobs



# Personal computers

1980s: Interactive jobs!




**UNIX®**





# 2000s: Cloud



- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.
- Hypervisor provides facilities to operating systems that OS provides to processes
  - multiplexes hardware among OS
  - protects and isolates OS from each other
- Hypervisors fundamentally enabled cloud computing

# 2000s: Smartphones

- New kinds of higher-level services: localisation, cellular, accelerometer, touch interface, etc.
- Resource constraints: Power management, UI system, etc.
- Even higher-level services: voice recognition, augmented reality, etc.
- Increased security concerns because of increase in sensitive data with rise of mobile banking, UPI etc. and because of moving devices







# **2000s: Cyber-physical systems**

**OS must not crash! Formally verified OS. Example: seL4**




# Typical progression of systems research

- Systems optimise for the “common case”
  - If common case changes, we need to rethink OS design
- Macro examples
  - Personal computers: batch jobs to interactive jobs
  - Smartphones: resource constraints, new sensors
  - Cyber physical systems: risk of human life

Users
Applications
Libraries
Operating systems
Hardware


# More trends: persistent memory

- Memory: volatile but fast
- Disk: persistent but large
- PMEM: persistent and fast!
  - PMEM aware file systems



# More trends: big memory

- OS were designed when memory was scarce: few KBs
- You can now buy a server with 12TB of DRAM!
- Transparent huge pages




# More trends: fast networks

- OS typically assumed network is *much slower* than DRAM
  - Far memory

2020s	Latency	Bandwidth
DRAM	15ns	400 GBps
Ethernet	500ns	50 GBps



# Rise of Unikernels



- OS optimises for common behaviours across all applications
- Each OS is now running only single application
- Unikernels optimise only for a single application

# Why should I care about learning OS?

If I don't want to do systems research

- OS is ultimately a study of abstraction. Absorb all the complexity away from the developer / the user.
- Principles are useful when designing any large-scale system