

Processes

Abhilash Jindal

Agenda

- Vocabulary (OSTEP Ch. 4)
 - What is a process? System calls? Scheduler? Address space?
- Memory management (OSTEP Ch. 13-17)
 - How to manage and isolate memory? What are memory APIs? How are they implemented?
- Processes in action (xv6 Ch. 3: system calls, x86 protection, trap handlers)
 - Process control block, user stack<>kernel stack, sys call handling
- Scheduling (xv6 Ch5: context switching, OSTEP Ch. 6-9)
 - Response time, throughput, fairness

Process is a running program

- Load program from disk to memory
 - Exactly how we loaded OS
- Give control to the process. Jump cs, eip

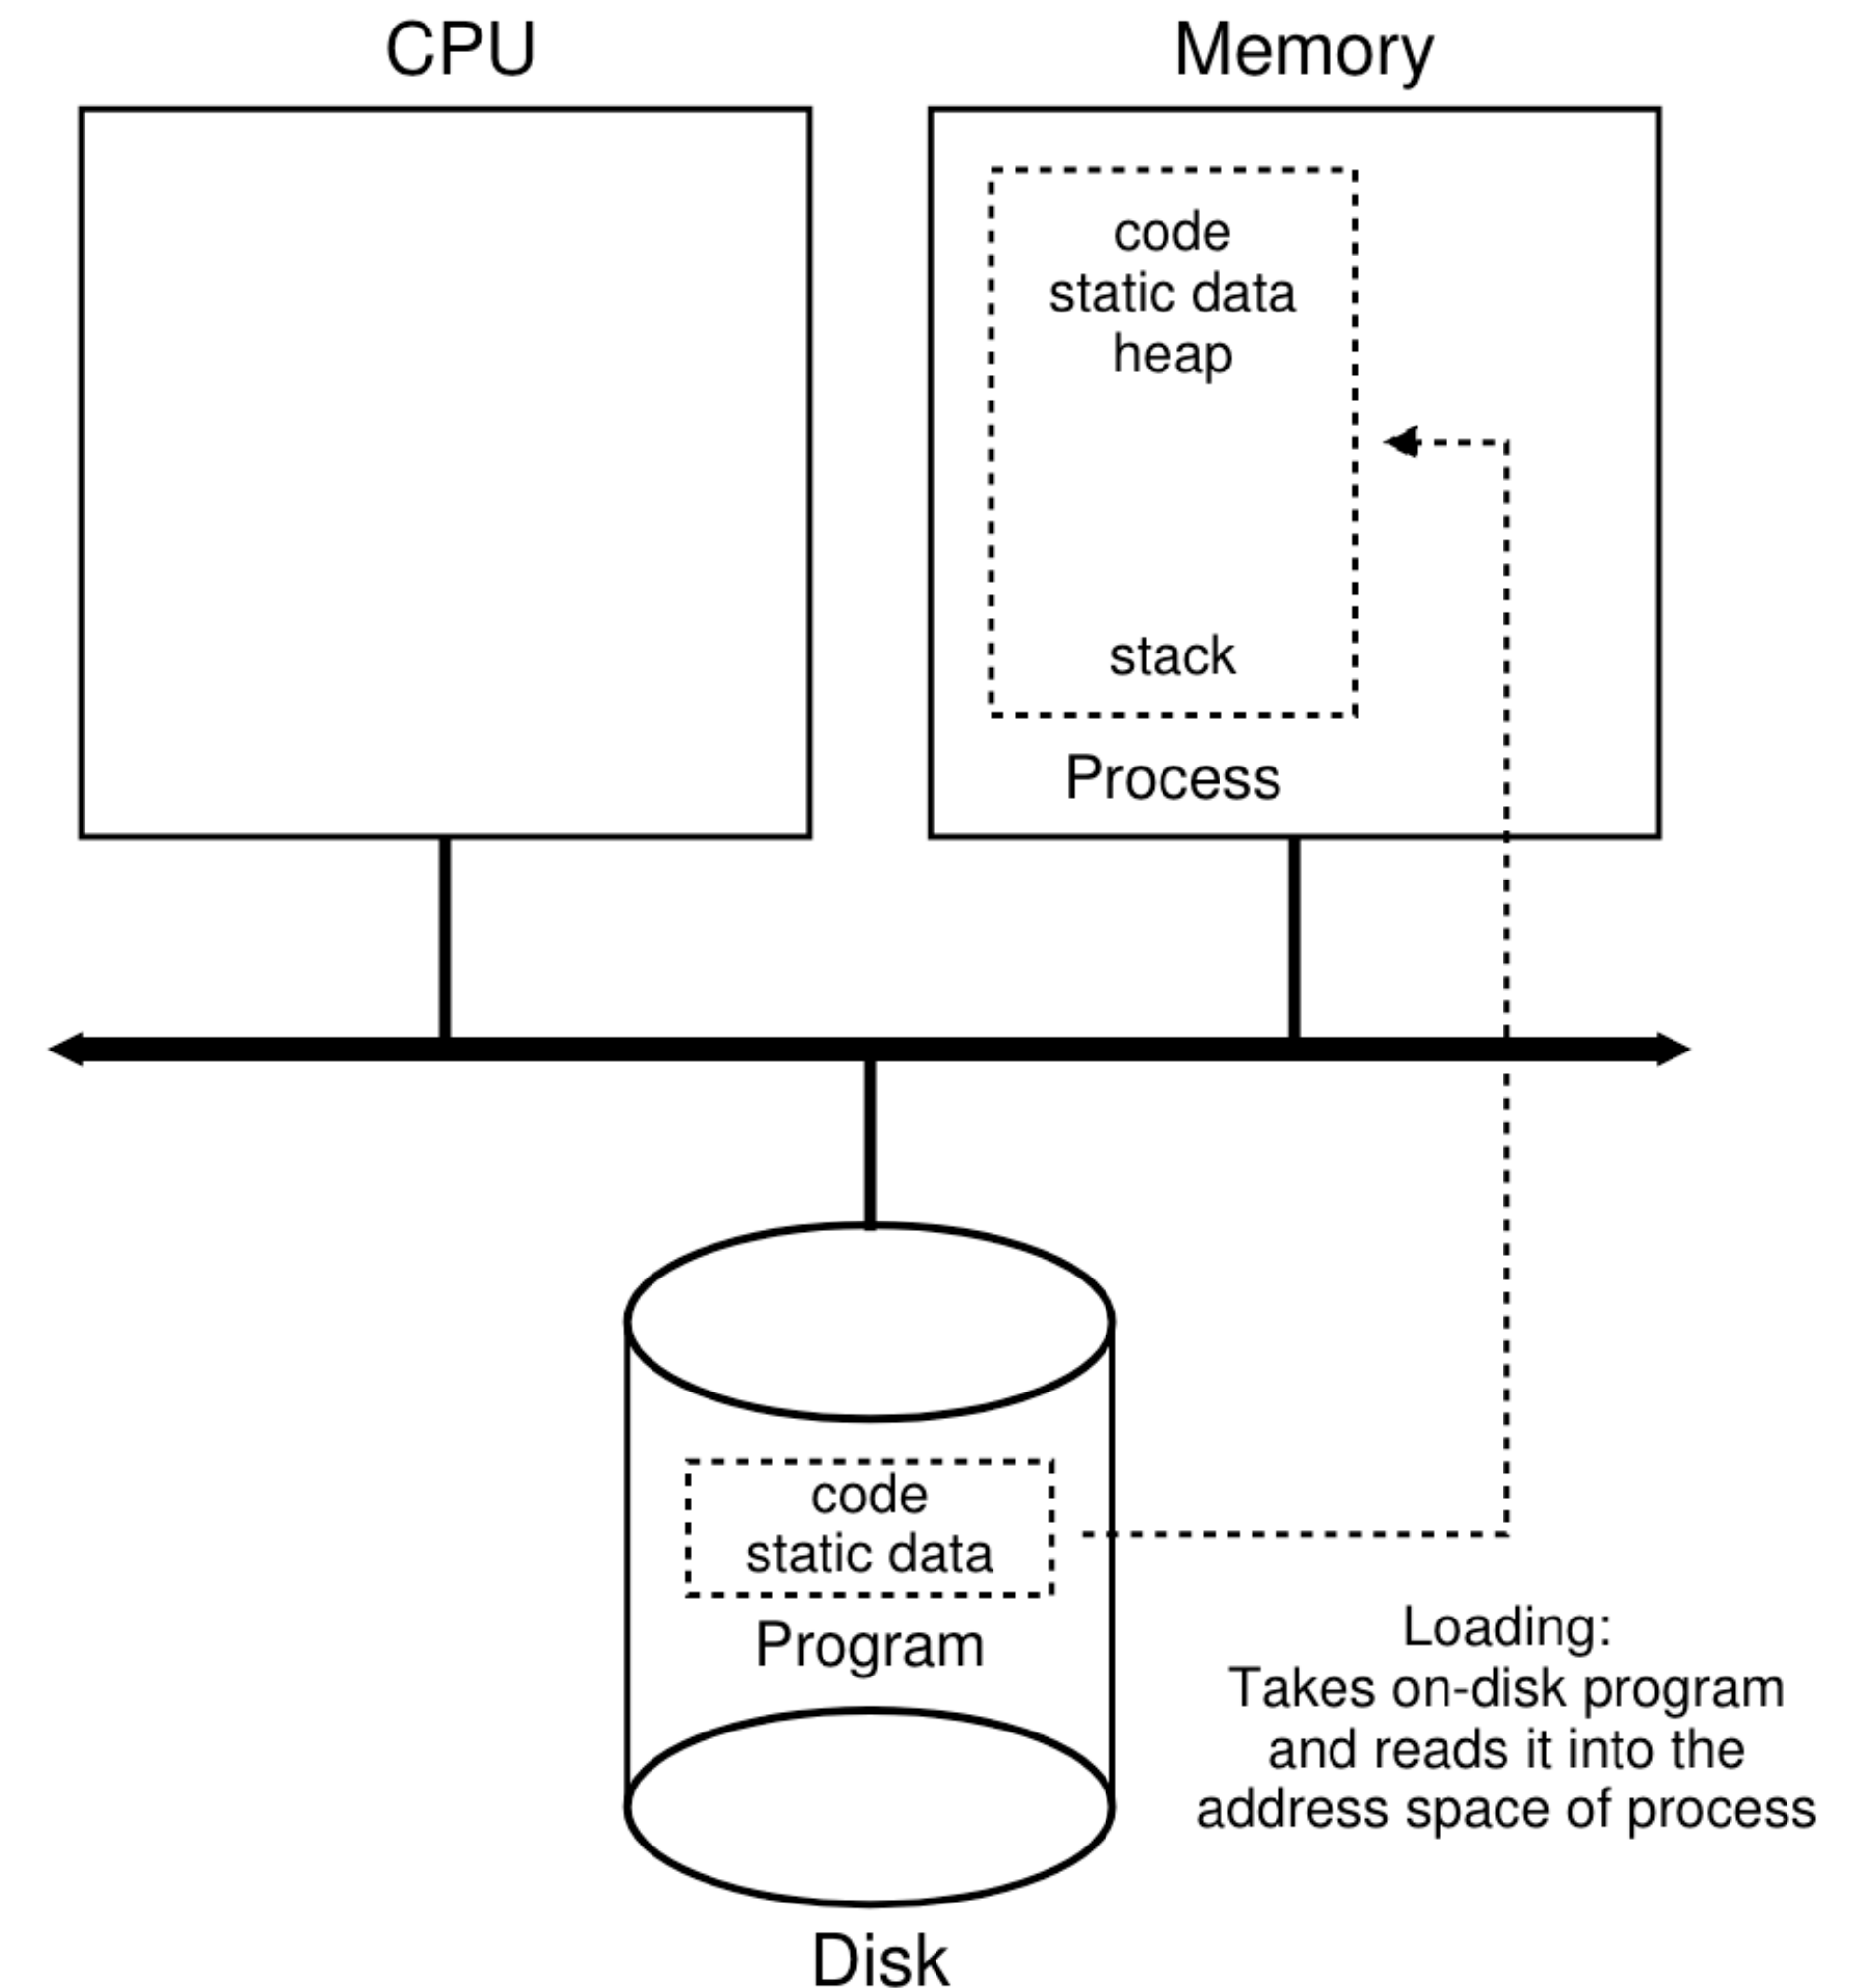


Figure 4.1: Loading: From Program To Process

Processes can ask OS to do work for them

System calls

```
$ strace cat /tmp/foo
...
openat(AT_FDCWD, "/tmp/foo", O_RDONLY) = 3
read(3, "hi\n", 131072) = 3
write(1, "hi\n", 3) = 3
...
```

OS maintains process states

Scheduler switches between processes

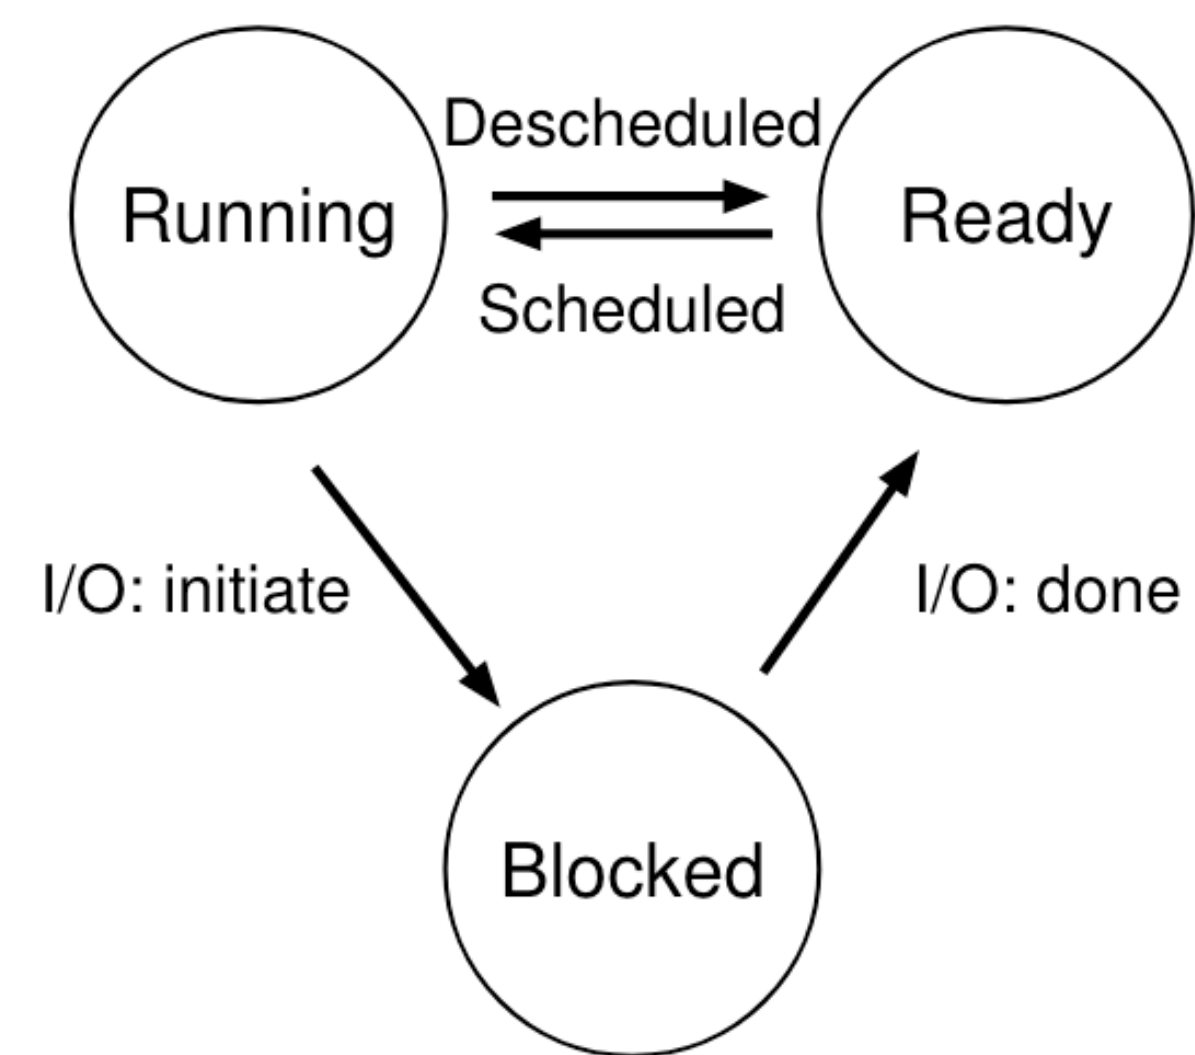


Figure 4.2: Process: State Transitions

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked,
5	Blocked	Running	so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

Figure 4.4: Tracing Process State: CPU and I/O

Calculator analogy: Computing long sum



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

Sharing the calculator



20
10
30
50
30
10
20
10



10
70
20
40
20
10
50
10

- Steps to share the calculator:
 - $20 + 10 = 30 + 30 = 60$
 - Write 60 in notebook, remember that we were done till 30, give calculator
 - $10 + 70 = 80$
 - Write 80 in notebook, remember that we were done till 70, give the calculator back


Memory isolation and management

OSTEP Ch. 13-17

Abhilash Jindal

Process Address Space

Code, Heap, Stack



Process address space

Function calling in action

Stack

02.s

eip

→

_foo:

pushl %ebp

movl %esp, %ebp

movl 8(%ebp), %eax

addl 12(%ebp), %eax

popl %ebp

retl

.globl _main

.p2align 4, 0x90

_main:

pushl %ebp

movl %esp, %ebp

subl \$24, %esp

movl \$0, -4(%ebp)

movl \$41, (%esp)

movl \$42, 4(%esp)

calll _foo

addl \$24, %esp

popl %ebp

retl

Save caller's base pointer

ebp = esp

eax = *(ebp + 8)

eax = eax + *(ebp + 12)

Restore caller's base pointer

change eip to return address

-- Begin function main

Save caller's base pointer

ebp = esp

esp = esp - 0x18

*(ebp-4)=0


*(esp) = 41

*(esp+4) = 42

Push current eip on to stack, jump to foo

esp = esp + 24 (Restore caller's esp)

Restore caller's ebp





Memory APIs and bugs

- malloc, free. va.c.
 - malloc is for dynamic allocation. Size is not known at compile time. Slower than stack allocations. Need to find free space.
- Null pointer dereference. null.c
- Memory leak. leak.c
- Buffer overflow. overflow.c
- Use after free. useafterfree.c
- Invalid free. invalidfree.c
- Double free. doublefree.c
- Uninitialised read. uninitread.c

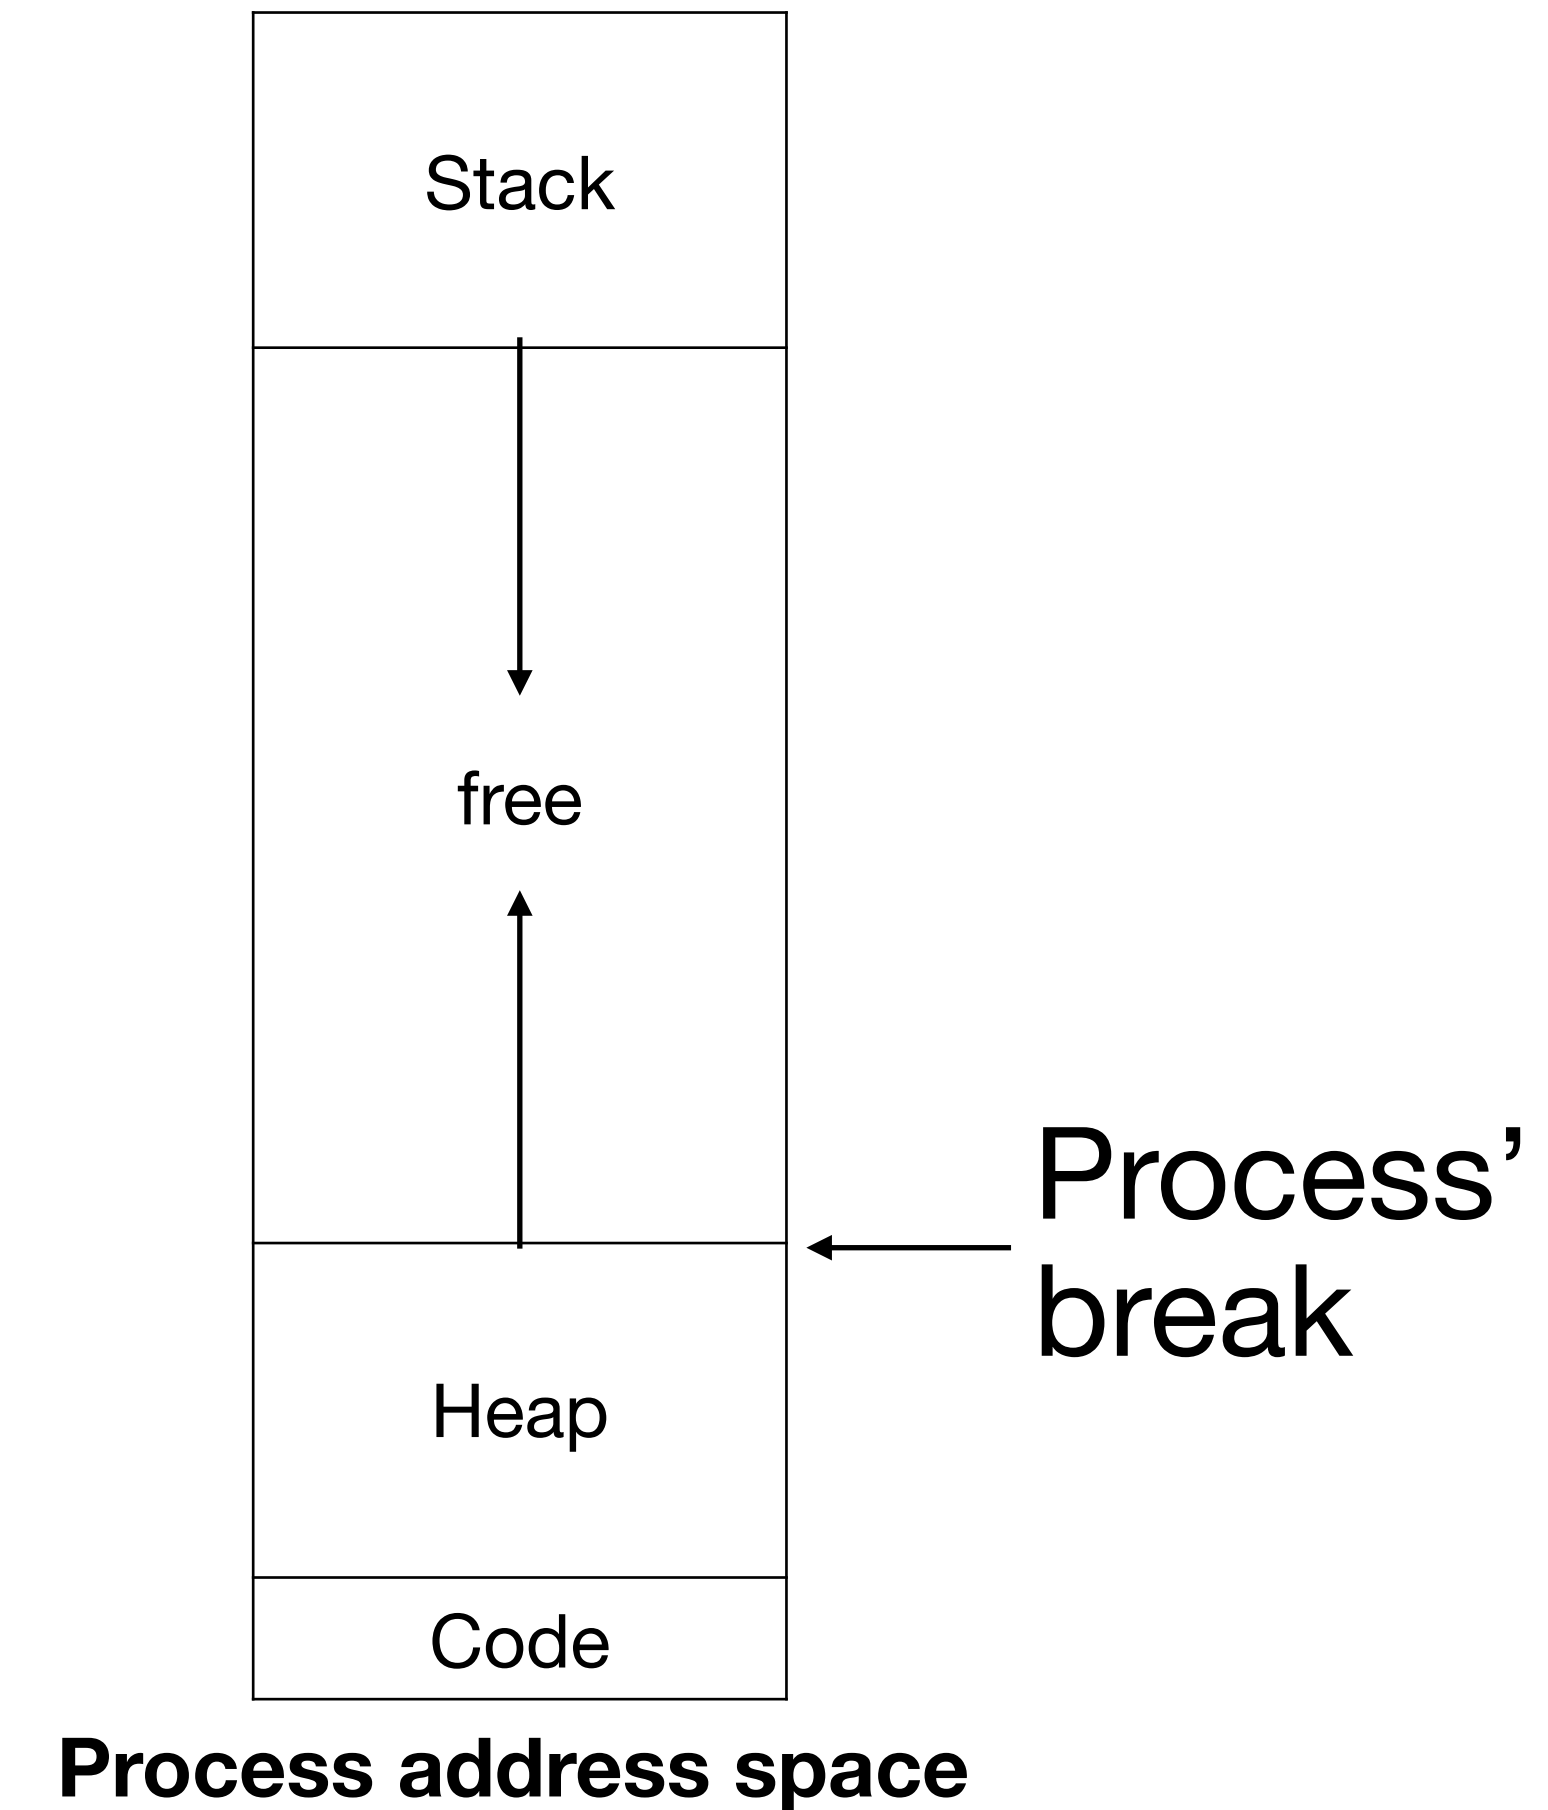
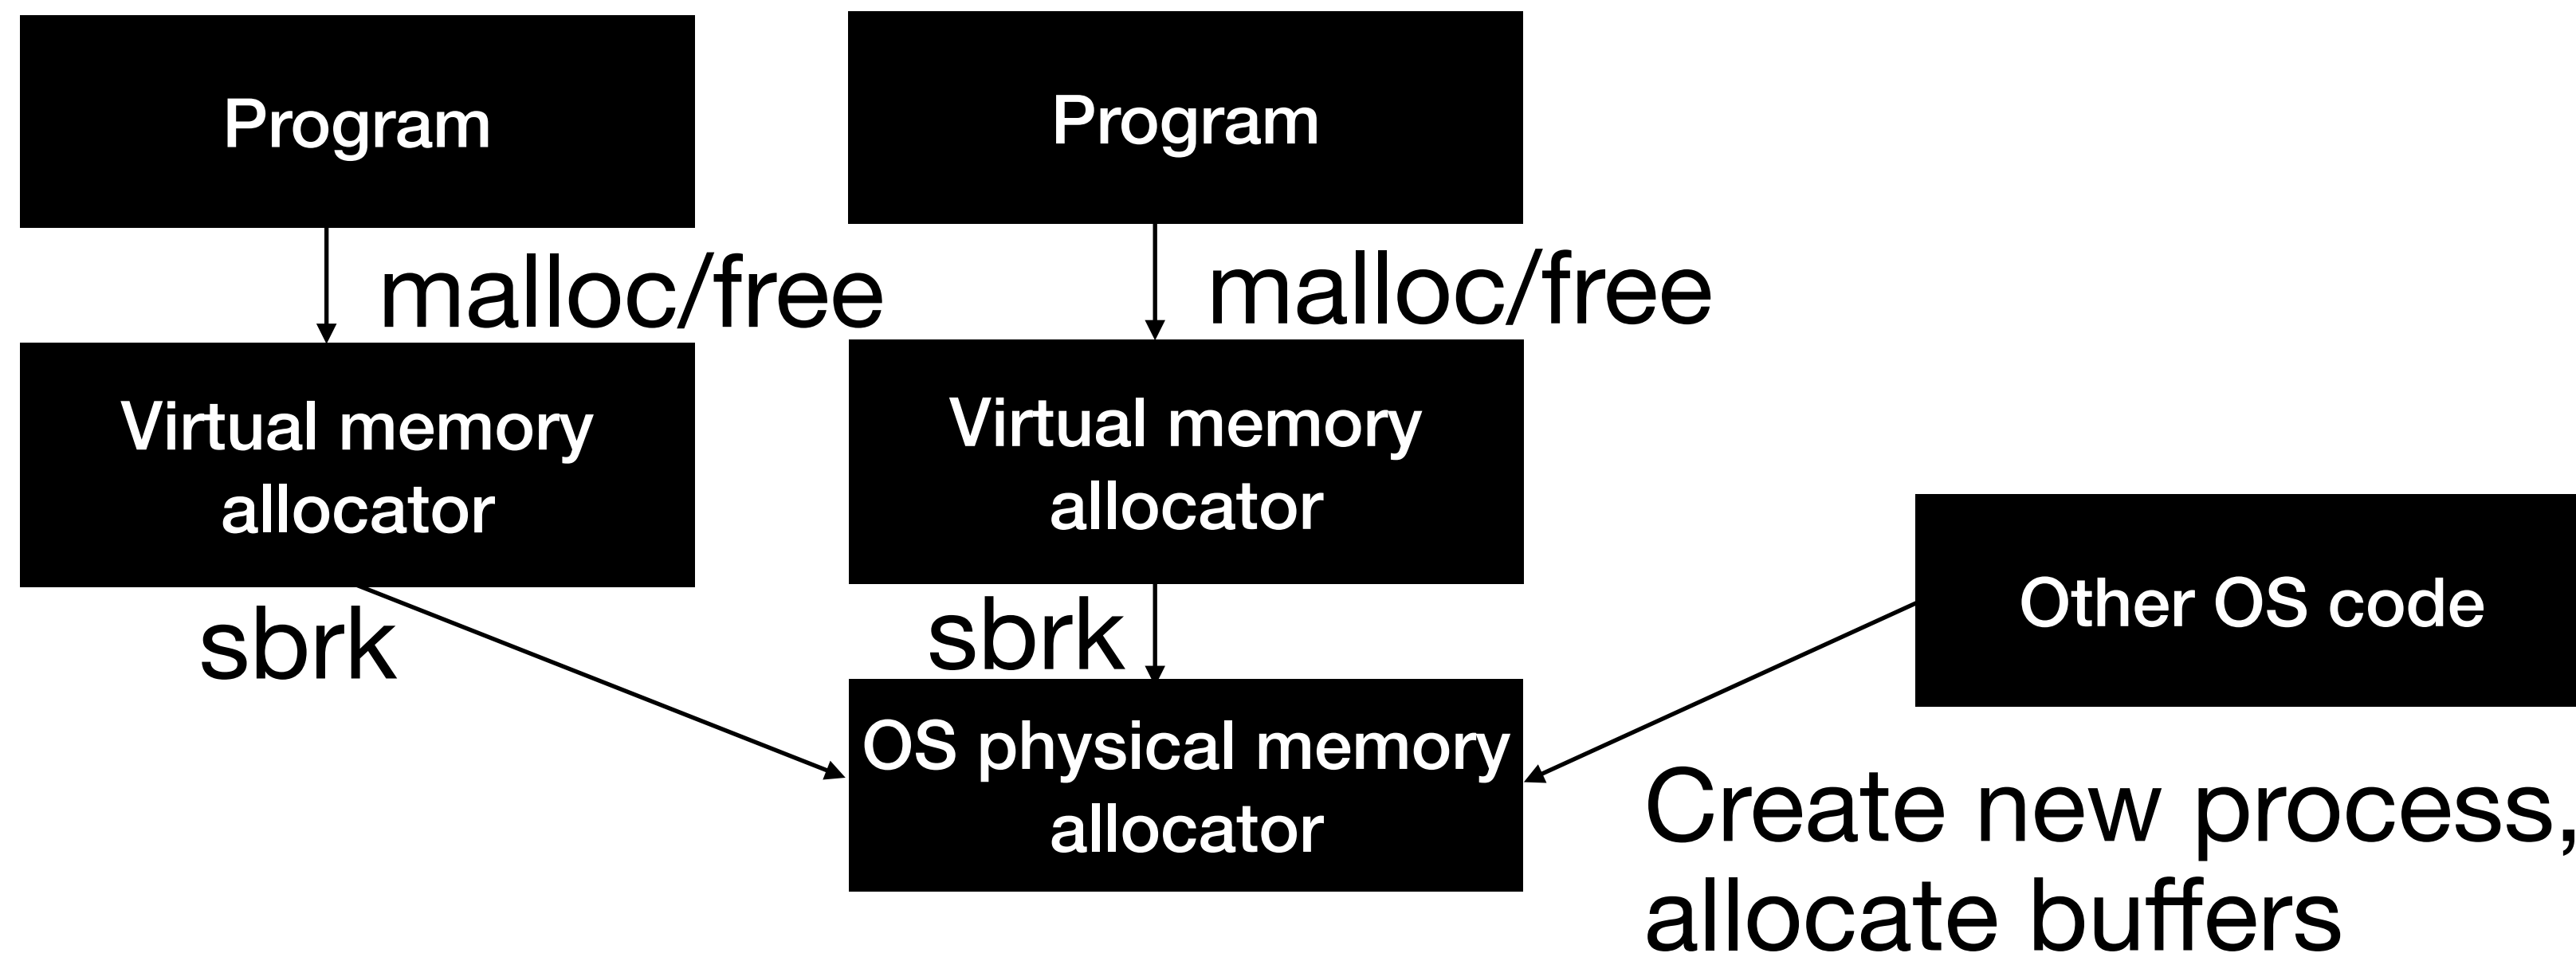
Memory allocator

Works with virtual memory



OS memory allocator

- `sbrk(int increment)` increments process' break. *increment* can be negative.



Memory allocation

```
ptr=malloc(size_t size);  
free(ptr);
```





Figure 17.2: Specific Contents Of The Header

Stack allocations are faster than heap allocations. No need to find space.

Memory allocator



Fragmented heap over time

- Assumptions
 - Do not apriori know allocation size and order
 - Cannot move memory once it is allocated. Program might have the pointer to it.
- Goals
 - Quickly satisfy variable-sized memory allocation requests. How to track free memory?
 - Minimize fragmentation

Memory (de)allocation patterns

- Small mallocs can be frequent. Large mallocs are usually infrequent.
 - After malloc, program will initialise the memory area.
- “Clustered deaths”: Objects allocated together die together.

Free list splitting and coalescing

```
ptr = malloc(100)
sptr = malloc(100)
optr = malloc(100)
free(sptr)
free(ptr)
free(optr)
```




Figure 17.3: A Heap With One Free Chunk

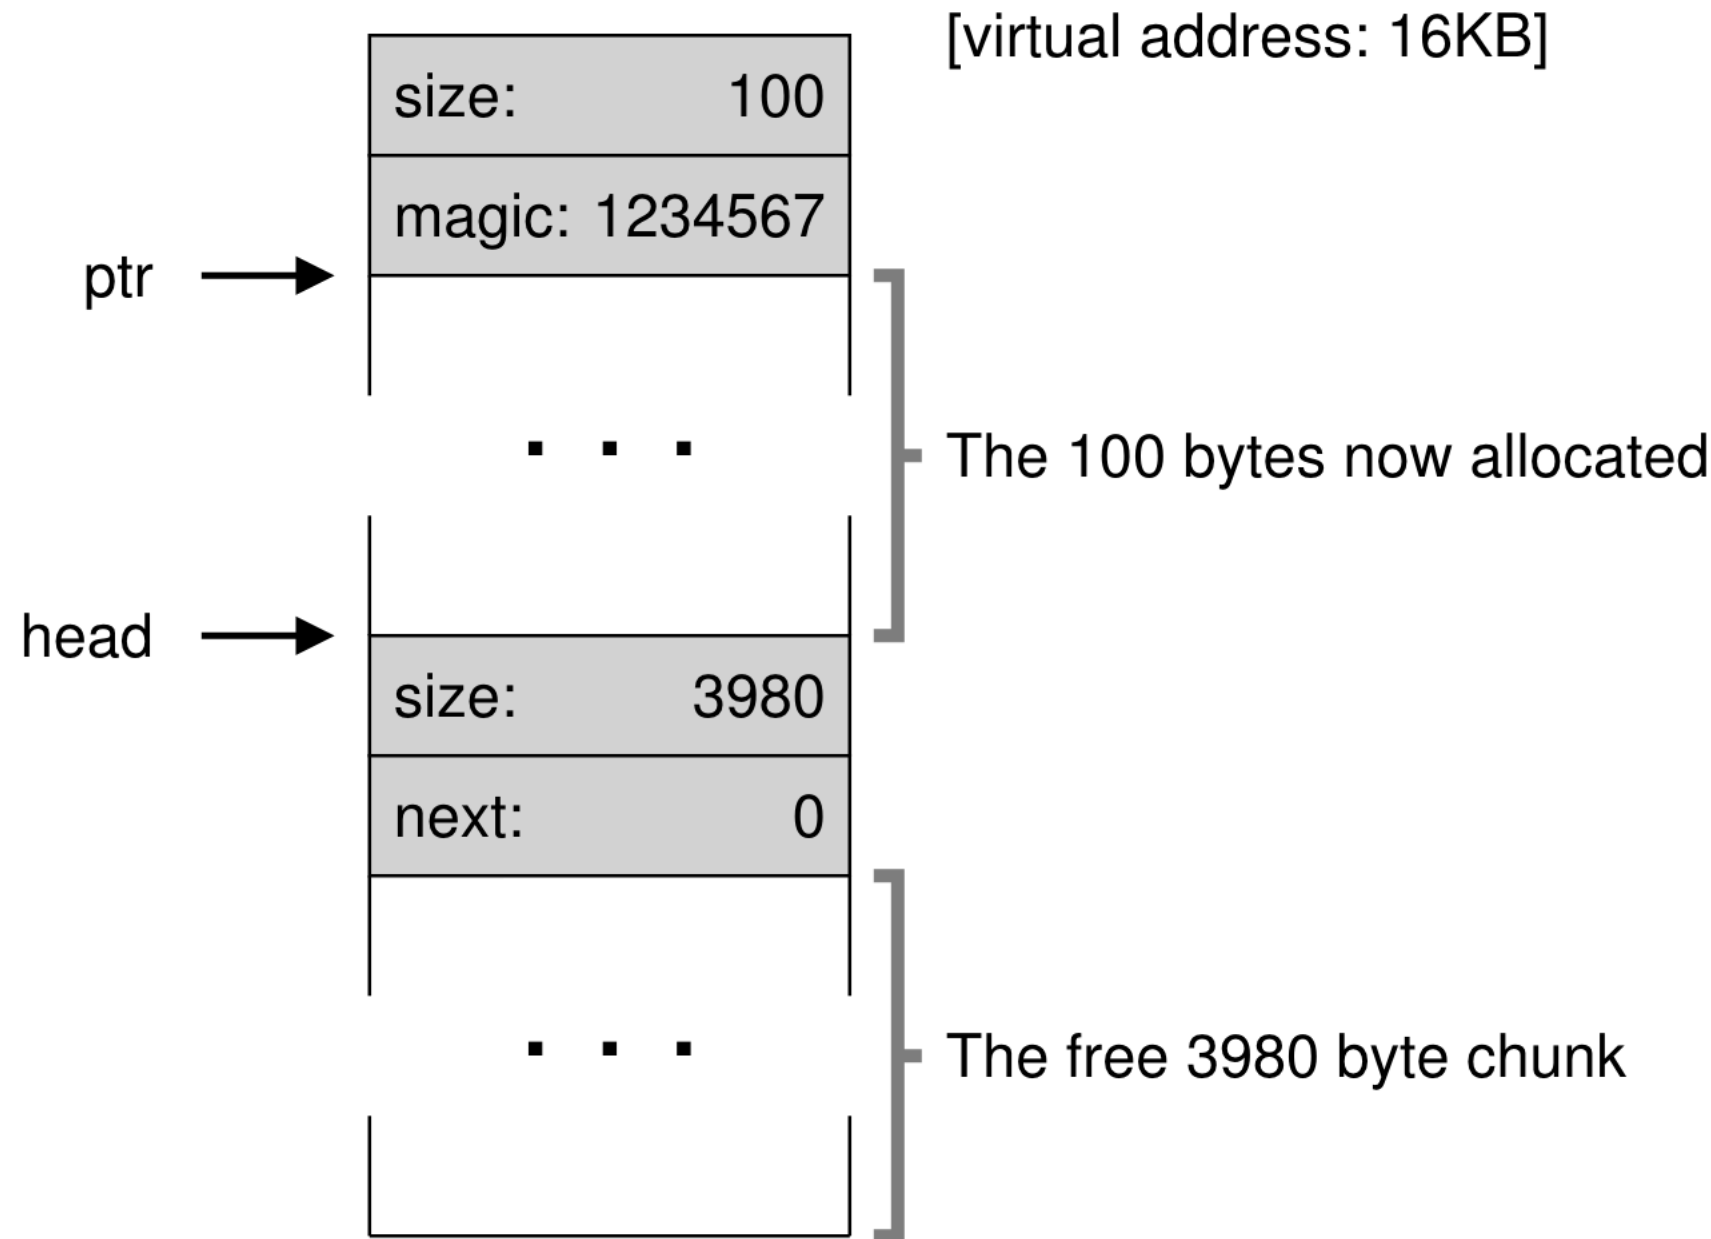


Figure 17.4: A Heap: After One Allocation

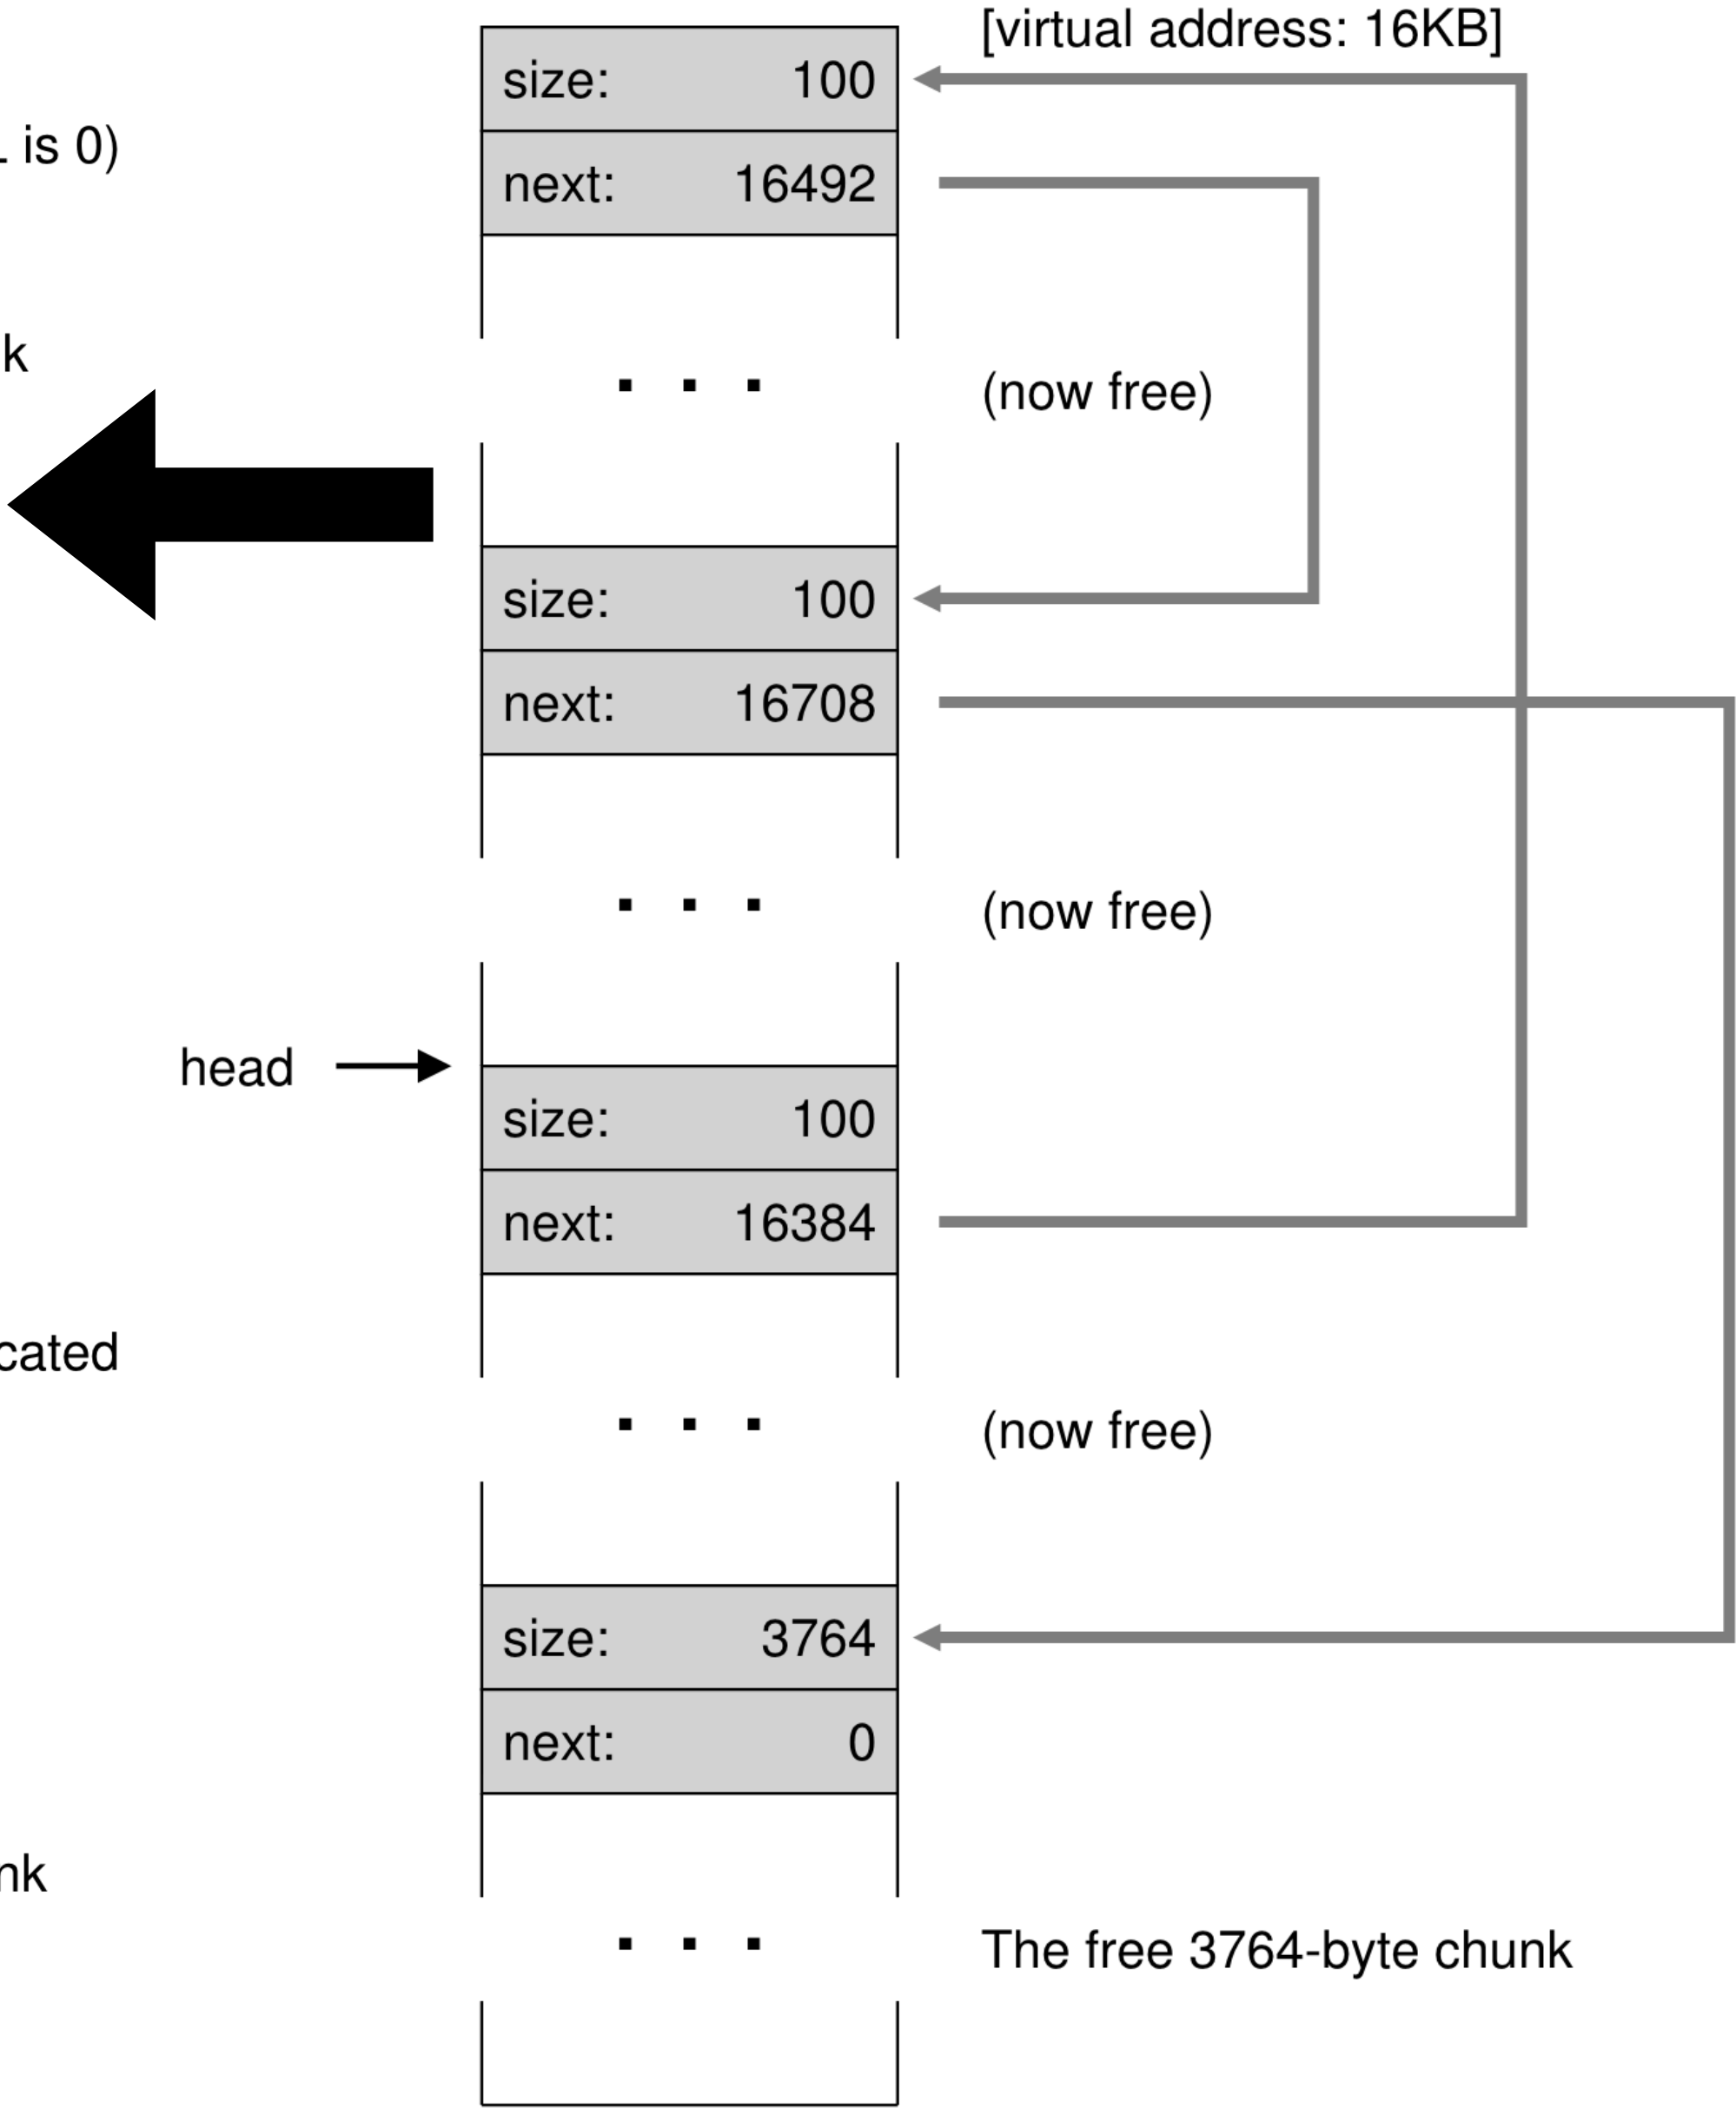


Figure 17.7: A Non-Coalesced Free List

Which block to allocate?



Example: malloc(15)

- Best fit
 - Slow. need to search the whole list
- First fit
 - Faster. (xv6: umalloc.c)
- Fragmentation
 - Example: malloc(25)



In which order to maintain lists?



- (De)allocation order
- Address order
 - Slow frees: need to traverse the free list
 - (xv6: umalloc.c)
 - Address order, first fit will allocate back-to-back allocations contiguously.
 - Due to “clustered deaths”, we may get better chances of coalescing




How to do coalescing?

Example: free(ptr)


- Straightforward in address order since we are traversing the free list in address order
- In deallocation order: when an area is freed, check if the “boundary tag” is present in the footer above
- First fit and address order work well in clustered deaths, simplify coalescing (no boundary tag), but cause fragmentation and traverse free list at free calls



First fit
address order



First fit
(de) allocation order



Boundary tag

