


Virtualization

Based on slides borrowed from Mythili, IIT Bombay

Virtualization terminology

- Hypervisor or virtual machine monitor (VMM): a piece of software that allows multiple VMs to run on a physical machine (PM)
- Guest OS runs inside the VM, and host OS runs on the PM




Why VMMs?

- Test and develop operating systems
- Run x86 OS on ARM hardware etc
- Backbone of cloud computing:
 - Renter perspective: I don't want to buy and manage physical machines.
 - Vendor perspective: I can overprovision and rent out a lot more virtual machines than available physical machines
 - Run competitor virtual machines with full isolation
 - Migrate virtual machines without downtime: load balance across physical machines, shut down physical machine for maintenance/power saving

Virtual machine monitor (VMM)


- Multiple Virtual Machines running on a Physical Machine – multiplex the underlying machine. Similar to how OS multiplexes processes on CPU



- VMM performs machine switch (much like context switch)
 - Run a VM for a bit, save context and switch to another VM, and so on...
- What is the problem?
 - Guest OS expects to have unrestricted access to hardware, runs privileged instructions, unlike user processes
 - But one guest cannot get access, must be isolated from other guests


Trap and emulate VMM (1)

- All CPUs have multiple privilege levels
 - Ring 0,1,2,3 in x86 CPUs
- Normally, user process in ring 3, OS in ring 0
 - Privileged instructions only run in ring 0
- Now, user process in ring 3, VMM/host OS in ring 0
 - Guest OS must be protected from guest apps
 - But not fully privileged like host OS/VMM
 - Can run in ring 1?
- Trap-and-emulate VMM: guest OS runs at lower privilege level than VMM, traps to VMM for privileged operation



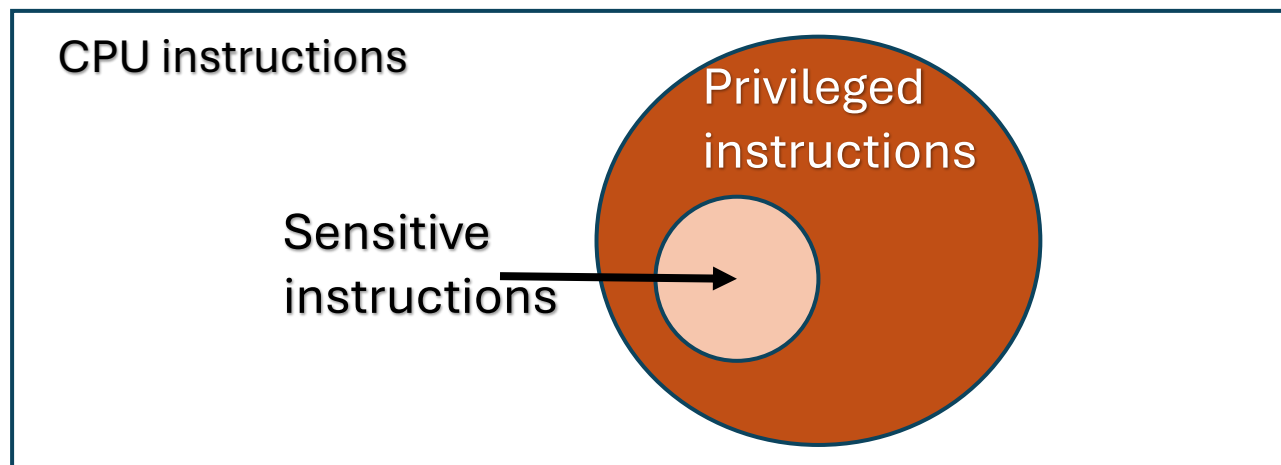
Trap and emulate VMM (2)

- Guest app has to handle syscall/interrupt
 - Special trap instr (int n), traps to VMM
 - VMM doesn't know how to handle trap
 - VMM jumps to guest OS trap handler
 - Trap handled by guest OS normally
- Guest OS performs return from trap
 - Privileged instructions trap to VMM
 - VMM jumps to corresponding user process
- Any privileged action by guest OS traps to VMM, emulated by VMM
 - Example: set IDT, set CR3, enable/disable interrupts
 - Sensitive data structures like IDT must be managed by VMM, not guest OS




Popek Goldberg theorem

- **Sensitive instruction** = changes hardware state, e.g. disable interrupts
- **Privileged instruction** = runs only in privileged mode
 - Traps to ring 0 if executed from unprivileged rings
- In order to build a VMM efficiently via trap-and-emulate method, **sensitive instructions should be a subset of privileged instructions**



Problems with trap and emulate

- Some x86 instructions which change hardware state (**sensitive instructions**) do not trap in less privileged ring 1!
- Why these problems?
 - OSes not developed to run at a lower privilege level
 - Instruction set architecture of x86 is not easily virtualizable (x86 wasn't designed with virtualization in mind)



x86 does not follow Popek-Goldberg theorem


Table 2.2: List of sensitive, unprivileged x86 instructions

Group	Instructions
Access to interrupt flag	pushf, popf, iret
Visibility into segment descriptors	lar, verr, verw, lsl
Segment manipulation instructions	pop <seg>, push <seg>, mov <seg>
Read-only access to privileged state	sgdt, sldt, sidt, smsw
Interrupt and gate instructions	fcall, longjump, retfar, str, int <n>

Robin et.al. USENIX Security, 2000

Example: Problems with trap and emulate

- **Eflags** register is a set of CPU flags
 - IF (interrupt flag) indicates if interrupts on/off
- Consider the **popf** instruction in x86
 - Pops values on top of stack and sets eflags
- Executed in ring 0, all flags set normally
- Executed in ring 1, only some flags set
 - IF is not set as it is privileged flag
- So, **popf is a sensitive instruction**, not privileged, **does not trap**, behaves differently when executed in different privilege levels
 - Guest OS is buggy in ring 1



Techniques to virtualize x86 (1)

- **Paravirtualization:** rewrite guest OS code to be virtualizable
 - Guest OS won't invoke privileged operations, makes “hypercalls” to VMM
 - Needs OS source code changes, cannot work with unmodified OS
 - Example: [Xen](#) hypervisor
- **Full virtualization:** CPU instructions of guest OS are translated to be virtualizable
 - Sensitive instructions translated to trap to VMM
 - Dynamic (on the fly) binary translation, so works with unmodified OS
 - Higher overhead than paravirtualization
 - Example: [VMWare workstation](#)

Dynamic Binary Translation-- Example

```
int isPrime(int a) {  
    for (int i = 2; i < a; i++) {  
        if (a % i == 0) return 0;  
    }  
    return 1;  
}
```

C program

89 f9 be 02 00 00 00 39 ce 7d ...

Binary representation

```
isPrime:  mov    %ecx, %edi ; %ecx = %edi (a)  
          mov    %esi, $2  ; i = 2  
          cmp    %esi, %ecx ; is i >= a?  
          jge    prime     ; jump if yes  
nexti:    mov    %eax, %ecx ; set %eax = a  
          cdq                     ; sign-extend  
          idiv   %esi         ; a % i  
          test   %edx, %edx ; is remainder zero?  
          jz     notPrime    ; jump if yes  
          inc    %esi         ; i++  
          cmp    %esi, %ecx ; is i >= a?  
          jl     nexti       ; jump if no  
prime:    mov    %eax, $1    ; return value in %eax  
          ret  
notPrime: xor    %eax, %eax ; %eax = 0  
          ret
```

Assembly instructions

Dynamic Binary Translation– Example (2)

```
isPrime:  mov    %ecx, %edi ; %ecx = %edi (a)
          mov    %esi, $2  ; i = 2
          cmp    %esi, %ecx ; is i >= a?
          jge    prime     ; jump if yes
nexti:    mov    %eax, %ecx ; set %eax = a
          cdq                     ; sign-extend
          idiv   %esi       ; a % i
          test   %edx, %edx ; is remainder zero?
          jz     notPrime   ; jump if yes
          inc    %esi       ; i++
          cmp    %esi, %ecx ; is i >= a?
          jl     nexti      ; jump if no
prime:    mov    %eax, $1   ; return value in %eax
          ret
notPrime: xor    %eax, %eax ; %eax = 0
          ret
```

Assembly instructions

isPrime(49)
Translator



```
isPrime':  mov %ecx, %edi ; IDENT
           mov %esi, $2
           cmp %esi, %ecx
           jge [takenAddr] ; JCC
           jmp [fallthrAddr]
```

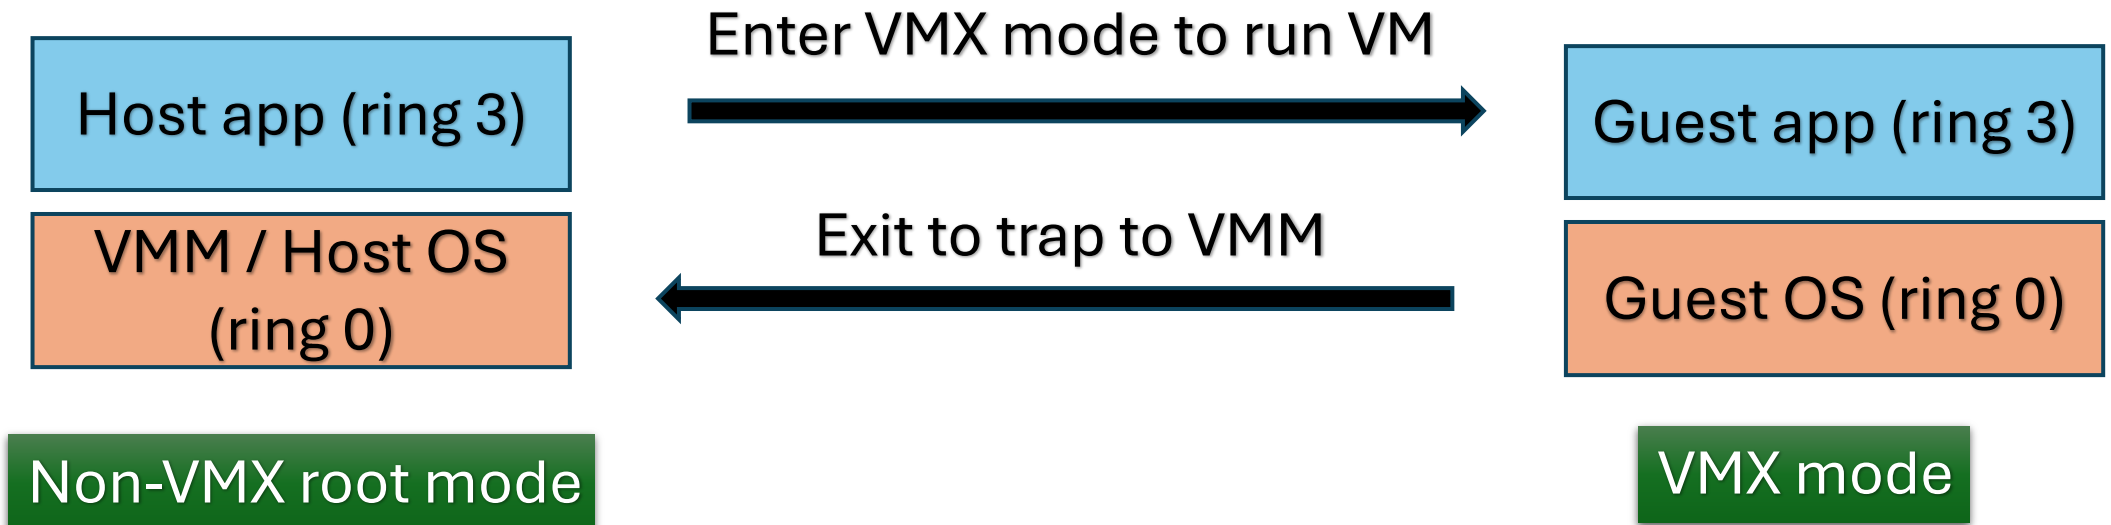
Compiled Code Fragment



```
isPrime':  *mov    %ecx, %edi ; IDENT
           mov    %esi, $2
           cmp    %esi, %ecx
           jge    [takenAddr] ; JCC
           ; fall-thru into next CCF
nexti':    *mov    %eax, %ecx ; IDENT
           cdq
           idiv   %esi
           test   %edx, %edx
           jz     notPrime' ; JCC
           ; fall-thru into next CCF
```

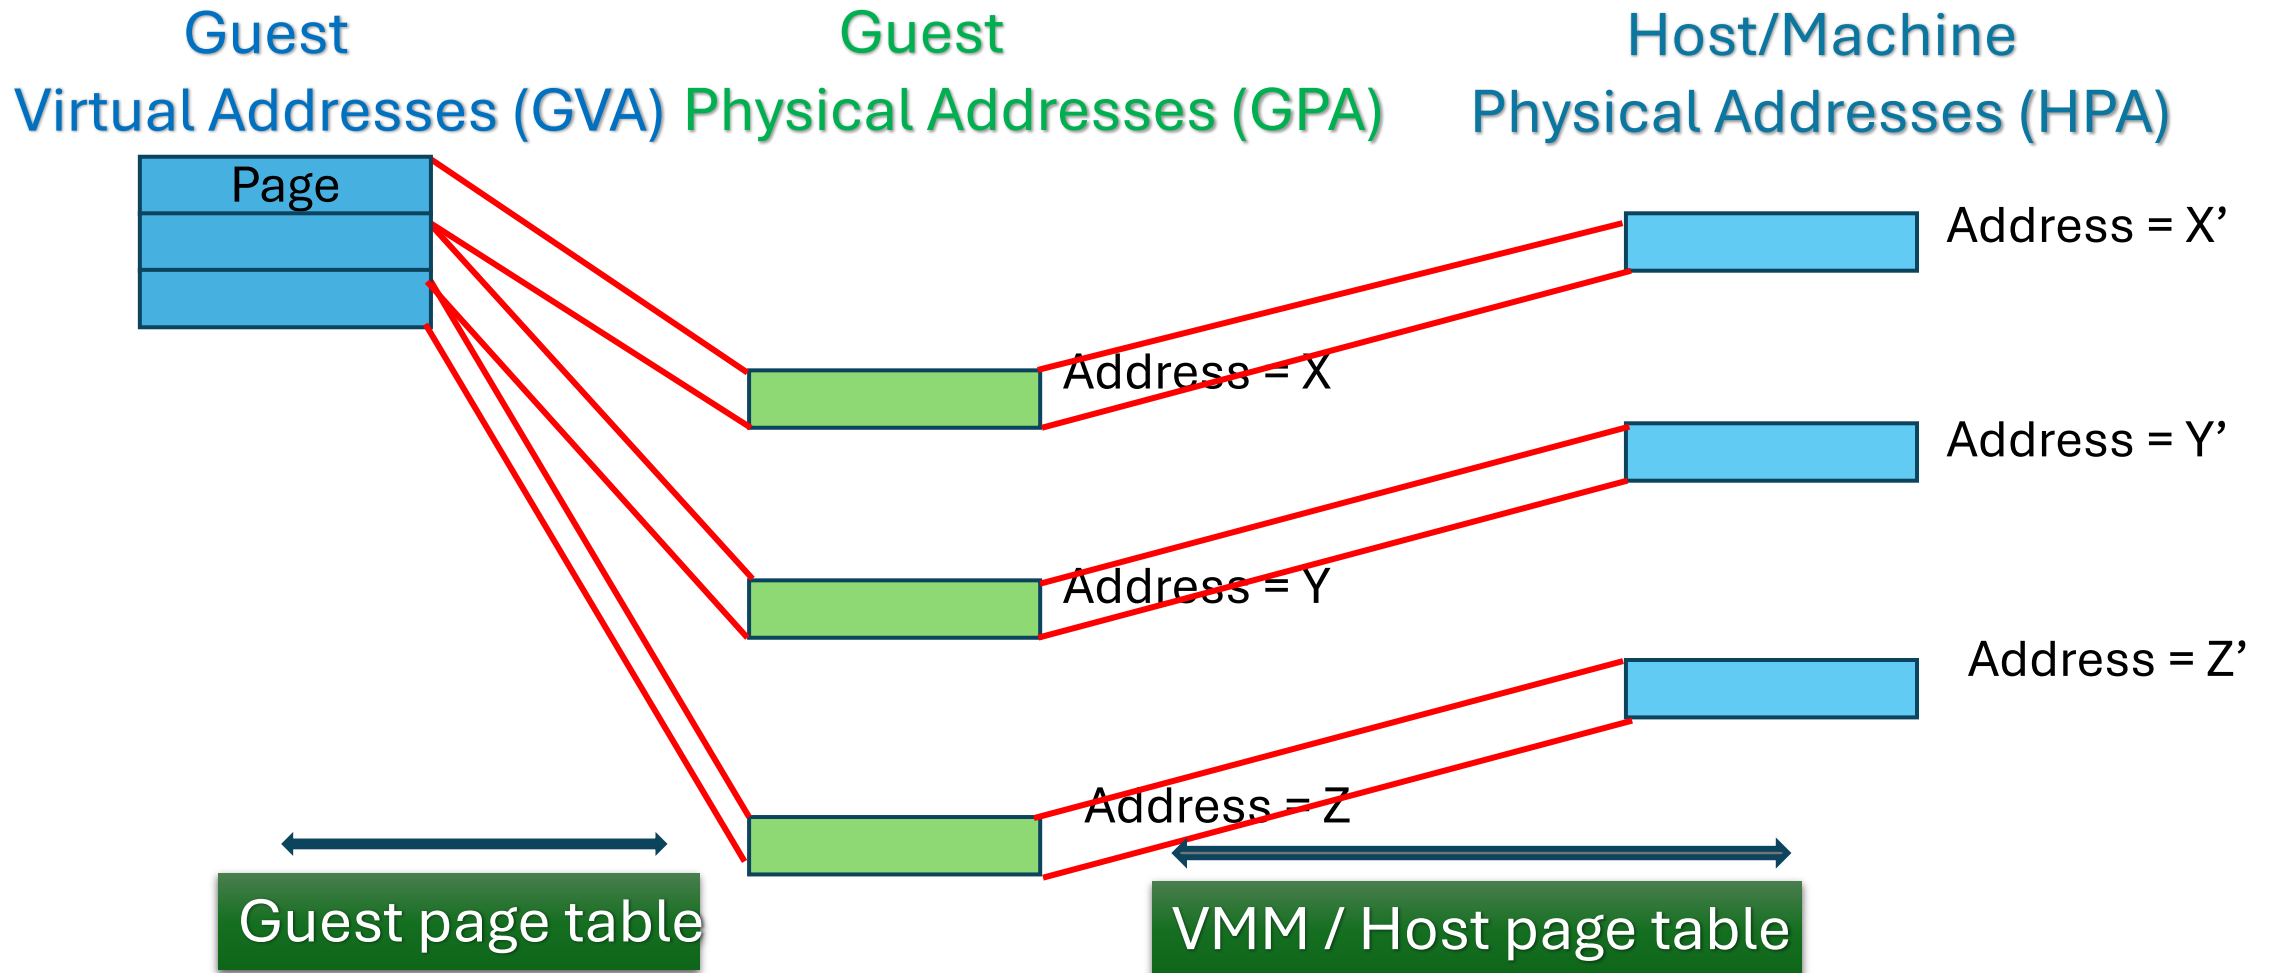
Techniques to virtualize x86 (2)

- **Hardware assisted virtualization:** KVM/QEMU in Linux
 - CPU has a special **VMX mode** of execution
 - X86 has 4 rings on non-VMX root mode, another 4 rings in VMX mode
- VMM enters VMX mode to run guest OS in (special) ring 0
- Exit back to VMM on triggers (VMM retains control)

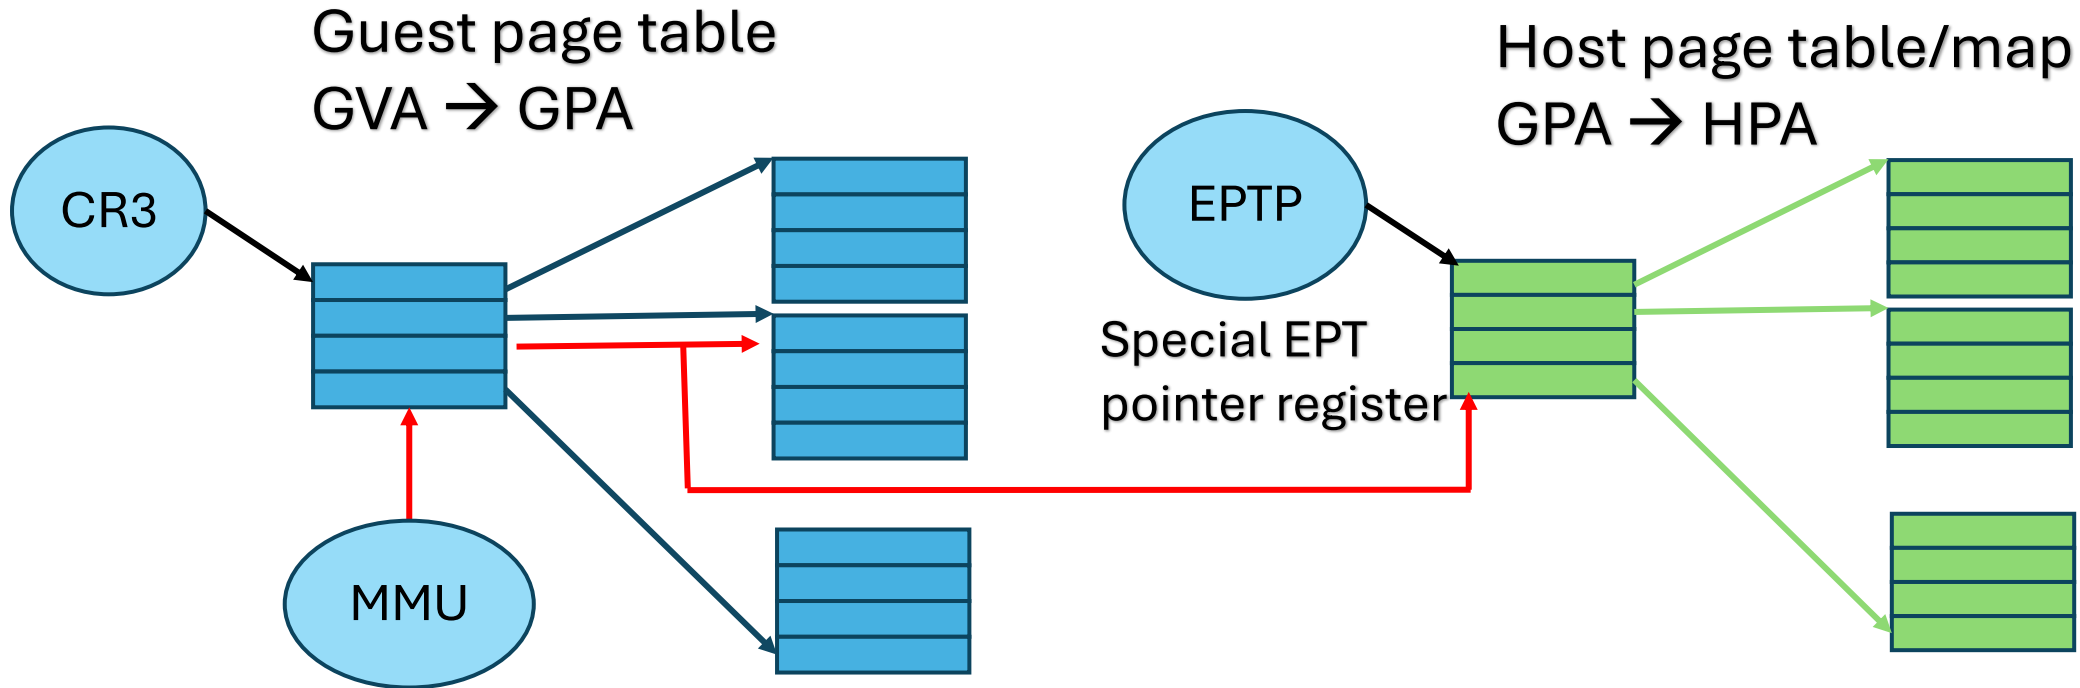


Memory virtualization

- What about address translation in virtual machines?



Extended page tables



- Page table walk by MMU: Start walking guest page table using GVA
- Guest PTE (for every level page table walk) gives GPA (cannot use GPA to access memory)
- Use GPA, walk host page table to find HPA, then access memory page, then next level access